PANDAS FOUNDATIONS

Visual exploratory data analysis

The iris data set

- Famous data set in pattern recognition

- 150 observations, 4 features each

  - Sepal length

  - Sepal width

  - Petal length

  - Petal width

- 3 species: setosa, versicolor, virginica

Data import

```
In [1]: import pandas as pd

In [2]: import matplotlib.pyplot as plt

In [3]: iris = pd.read_csv('iris.csv', index_col=0)

In [4]: print(iris.shape)
(150, 5)
```

Line plot

```
In [5]: iris.head()
Out[5]:
   sepal_length  sepal_width  petal_length  petal_width species
0           5.1          3.5           1.4          0.2  setosa
1           4.9          3.0           1.4          0.2  setosa
2           4.7          3.2           1.3          0.2  setosa
3           4.6          3.1           1.5          0.2  setosa
4           5.0          3.6           1.4          0.2  setosa

In [6]: iris.plot(x='sepal_length', y='sepal_width')

In [7]: plt.show()
```
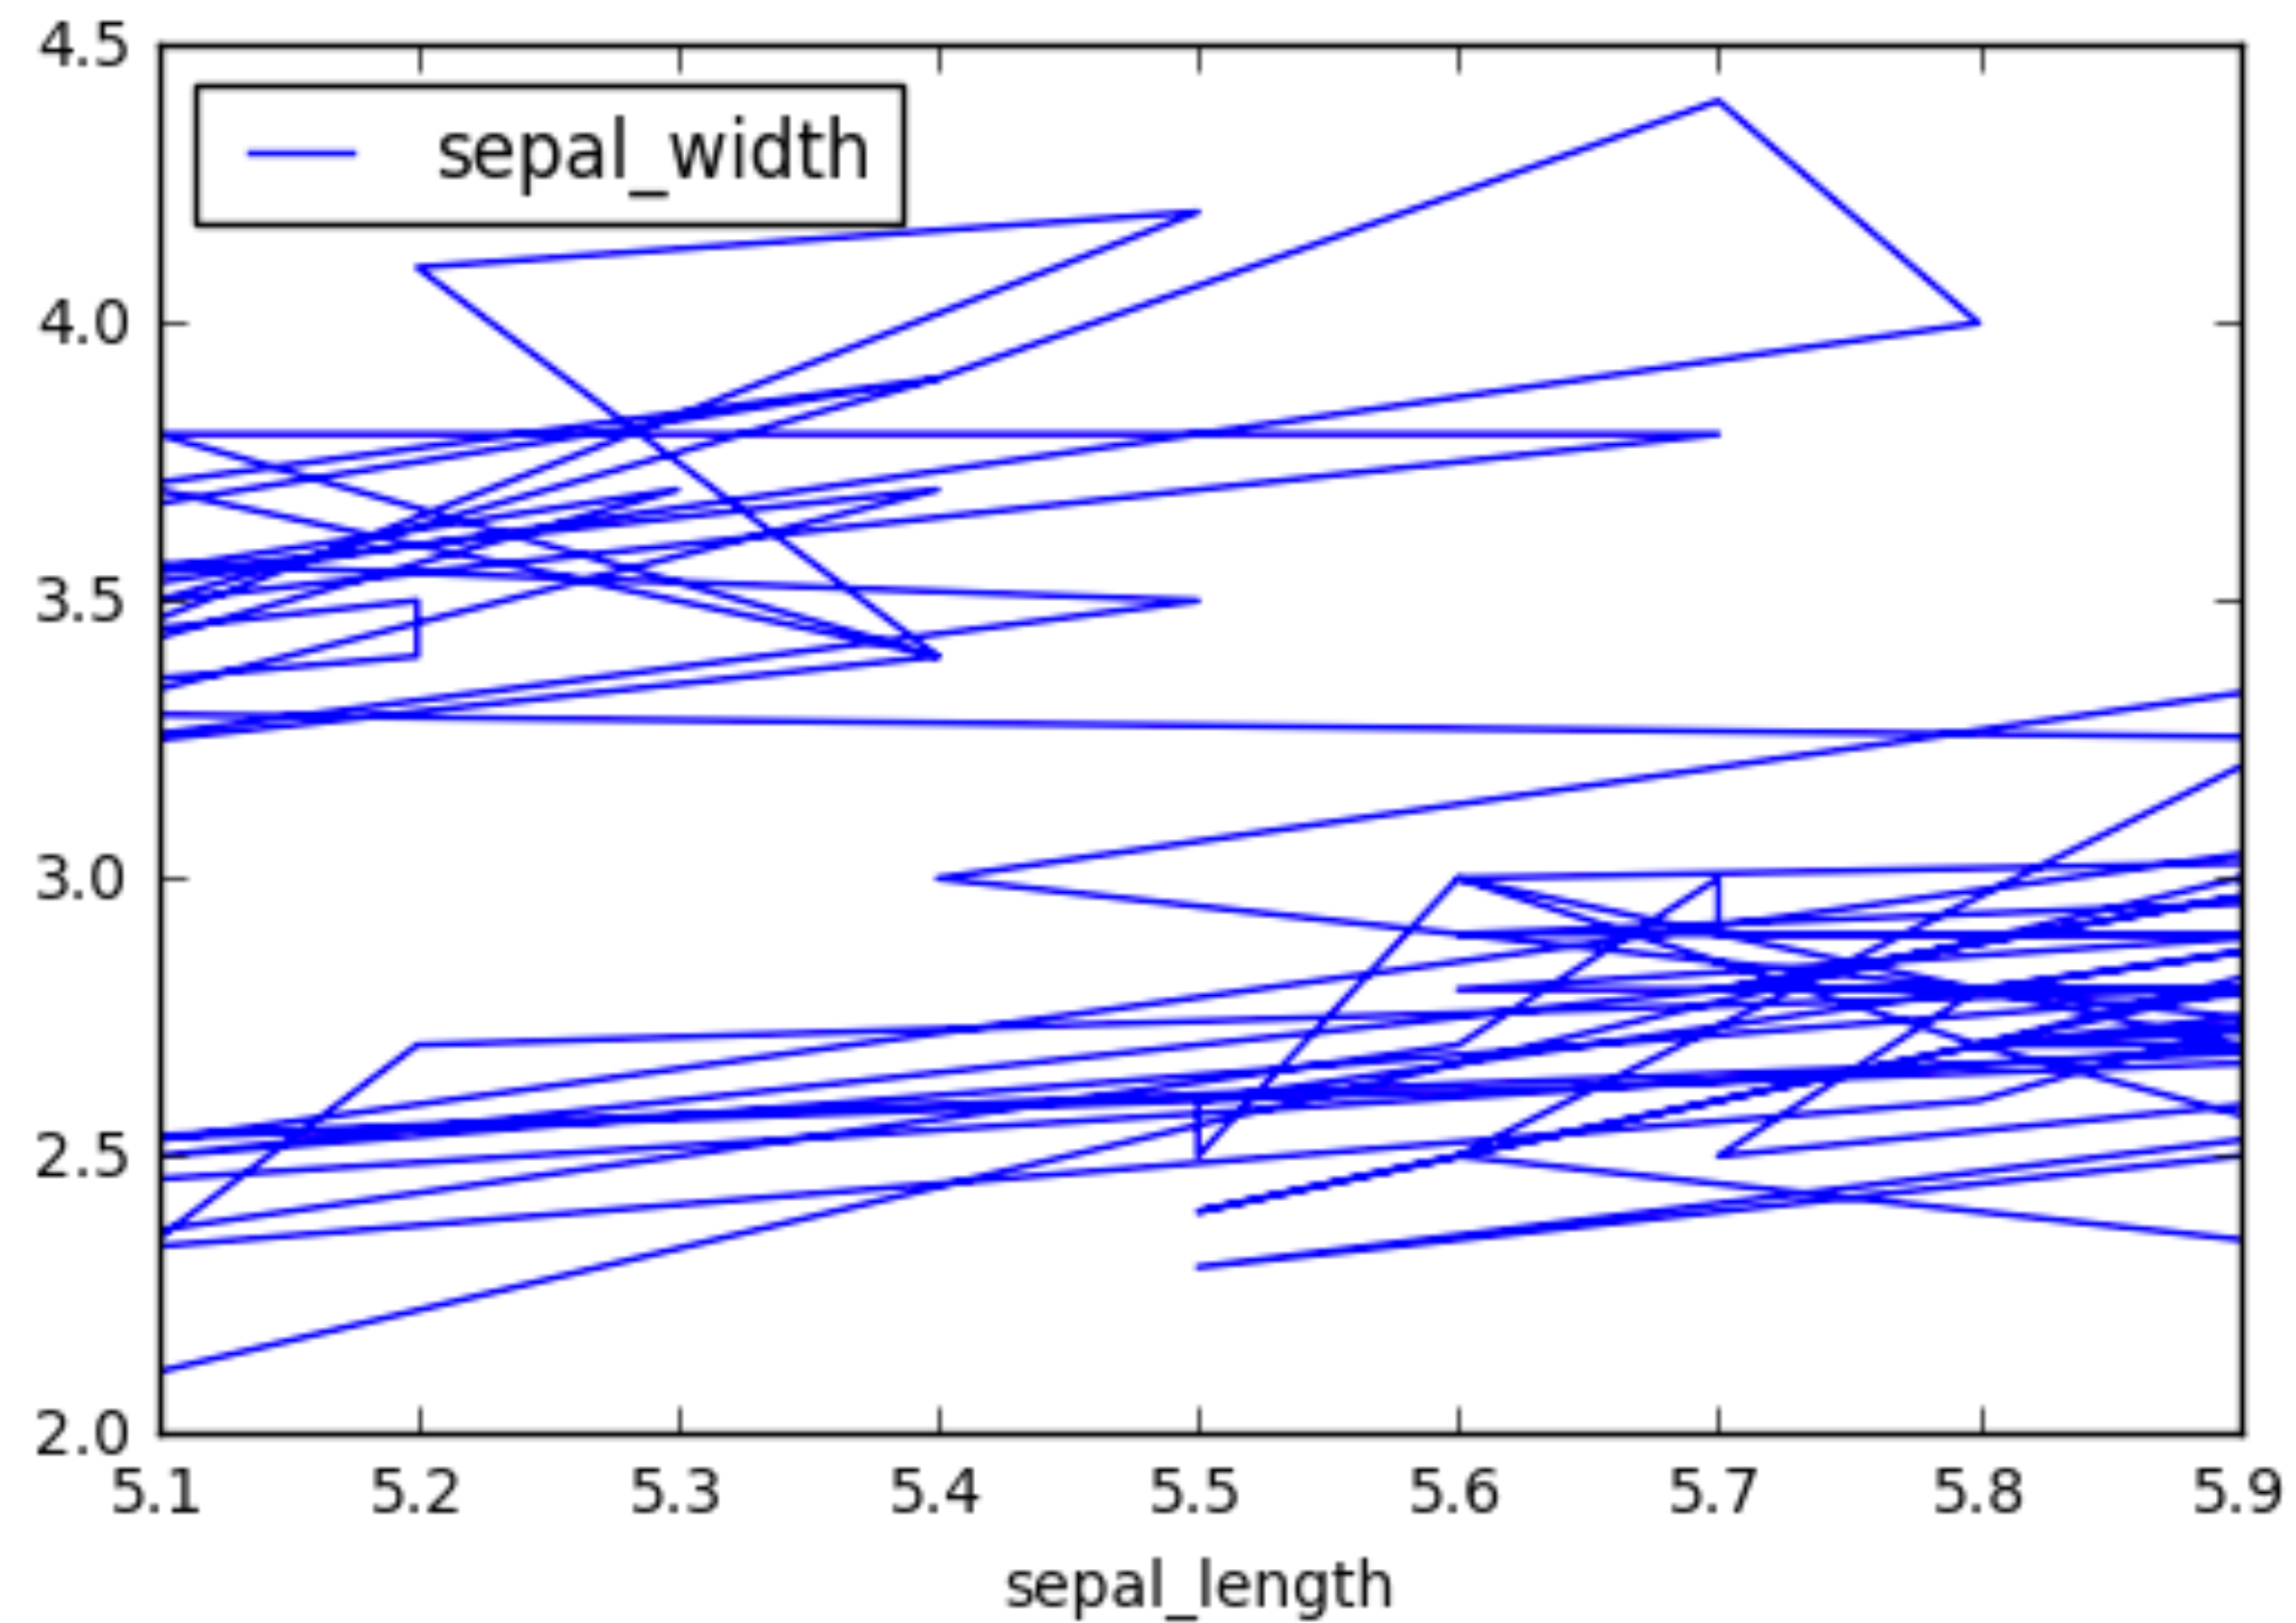
Line plot

Scatter plot

```
In [8]: iris.plot(x='sepal_length', y='sepal_width',
   ...:                kind='scatter')

In [9]: plt.xlabel('sepal length (cm)')

In [10]: plt.ylabel('sepal width (cm)')

In [11]: plt.show()
```
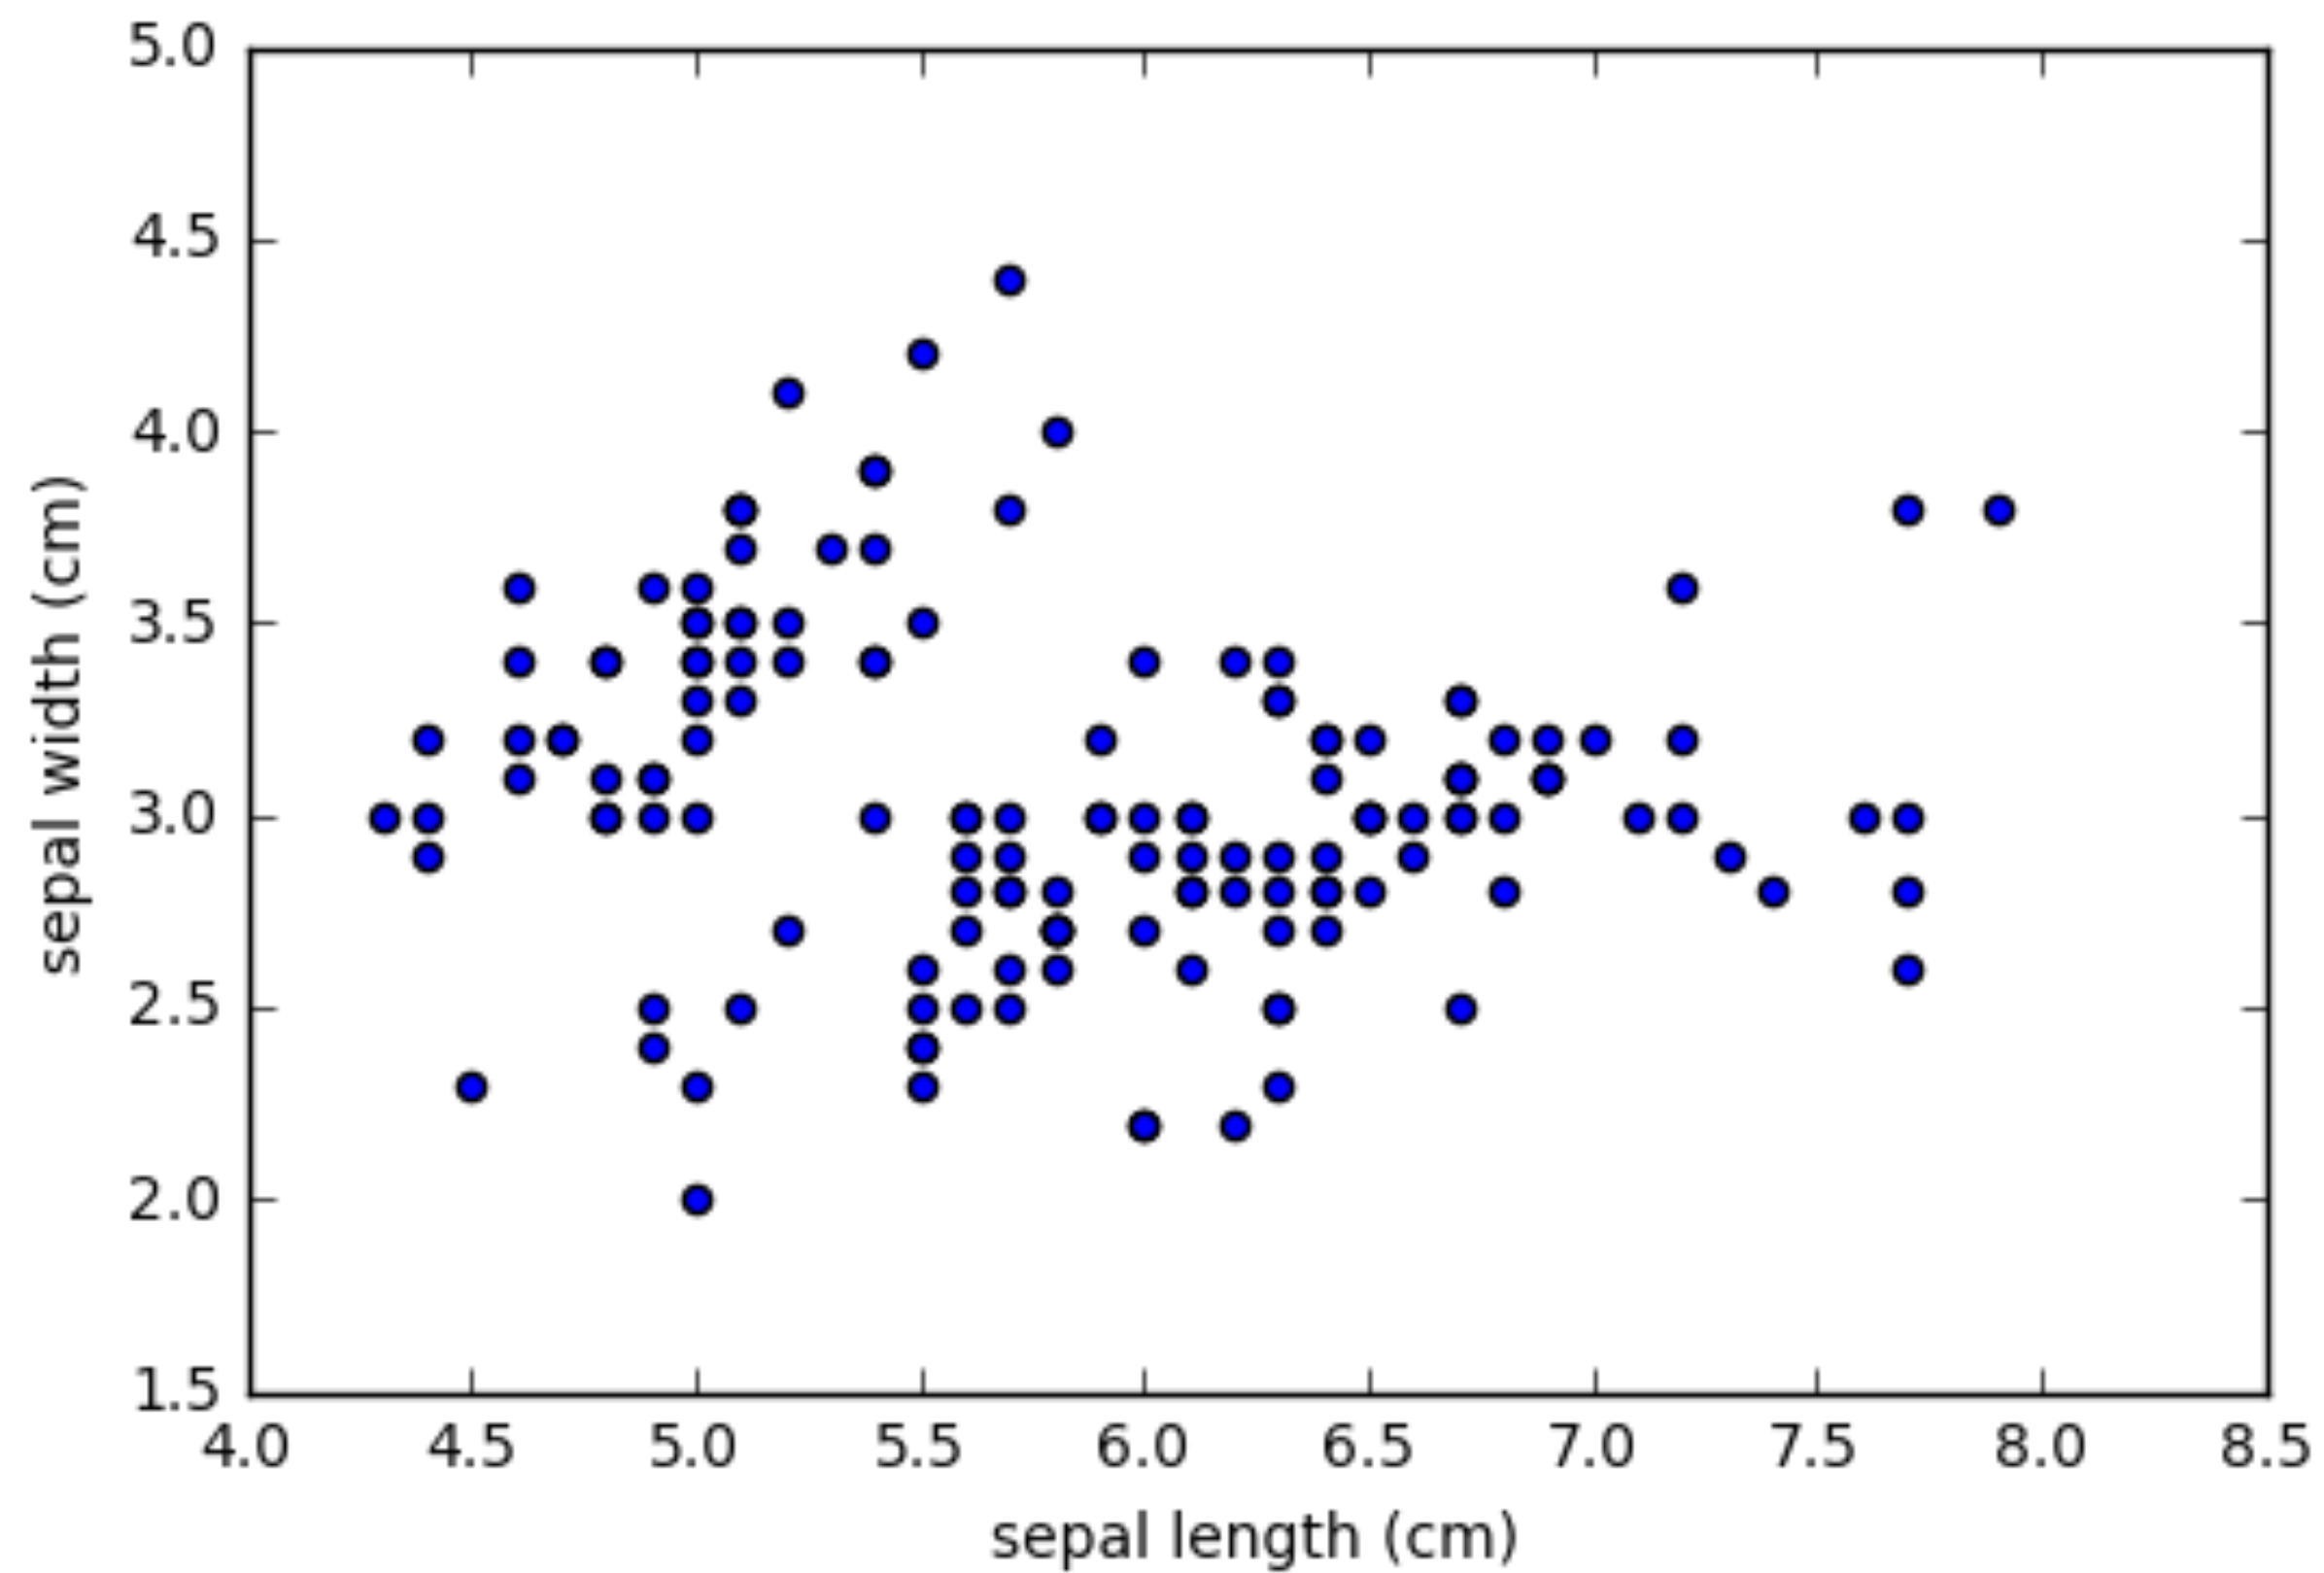
Scatter plot

Box plot
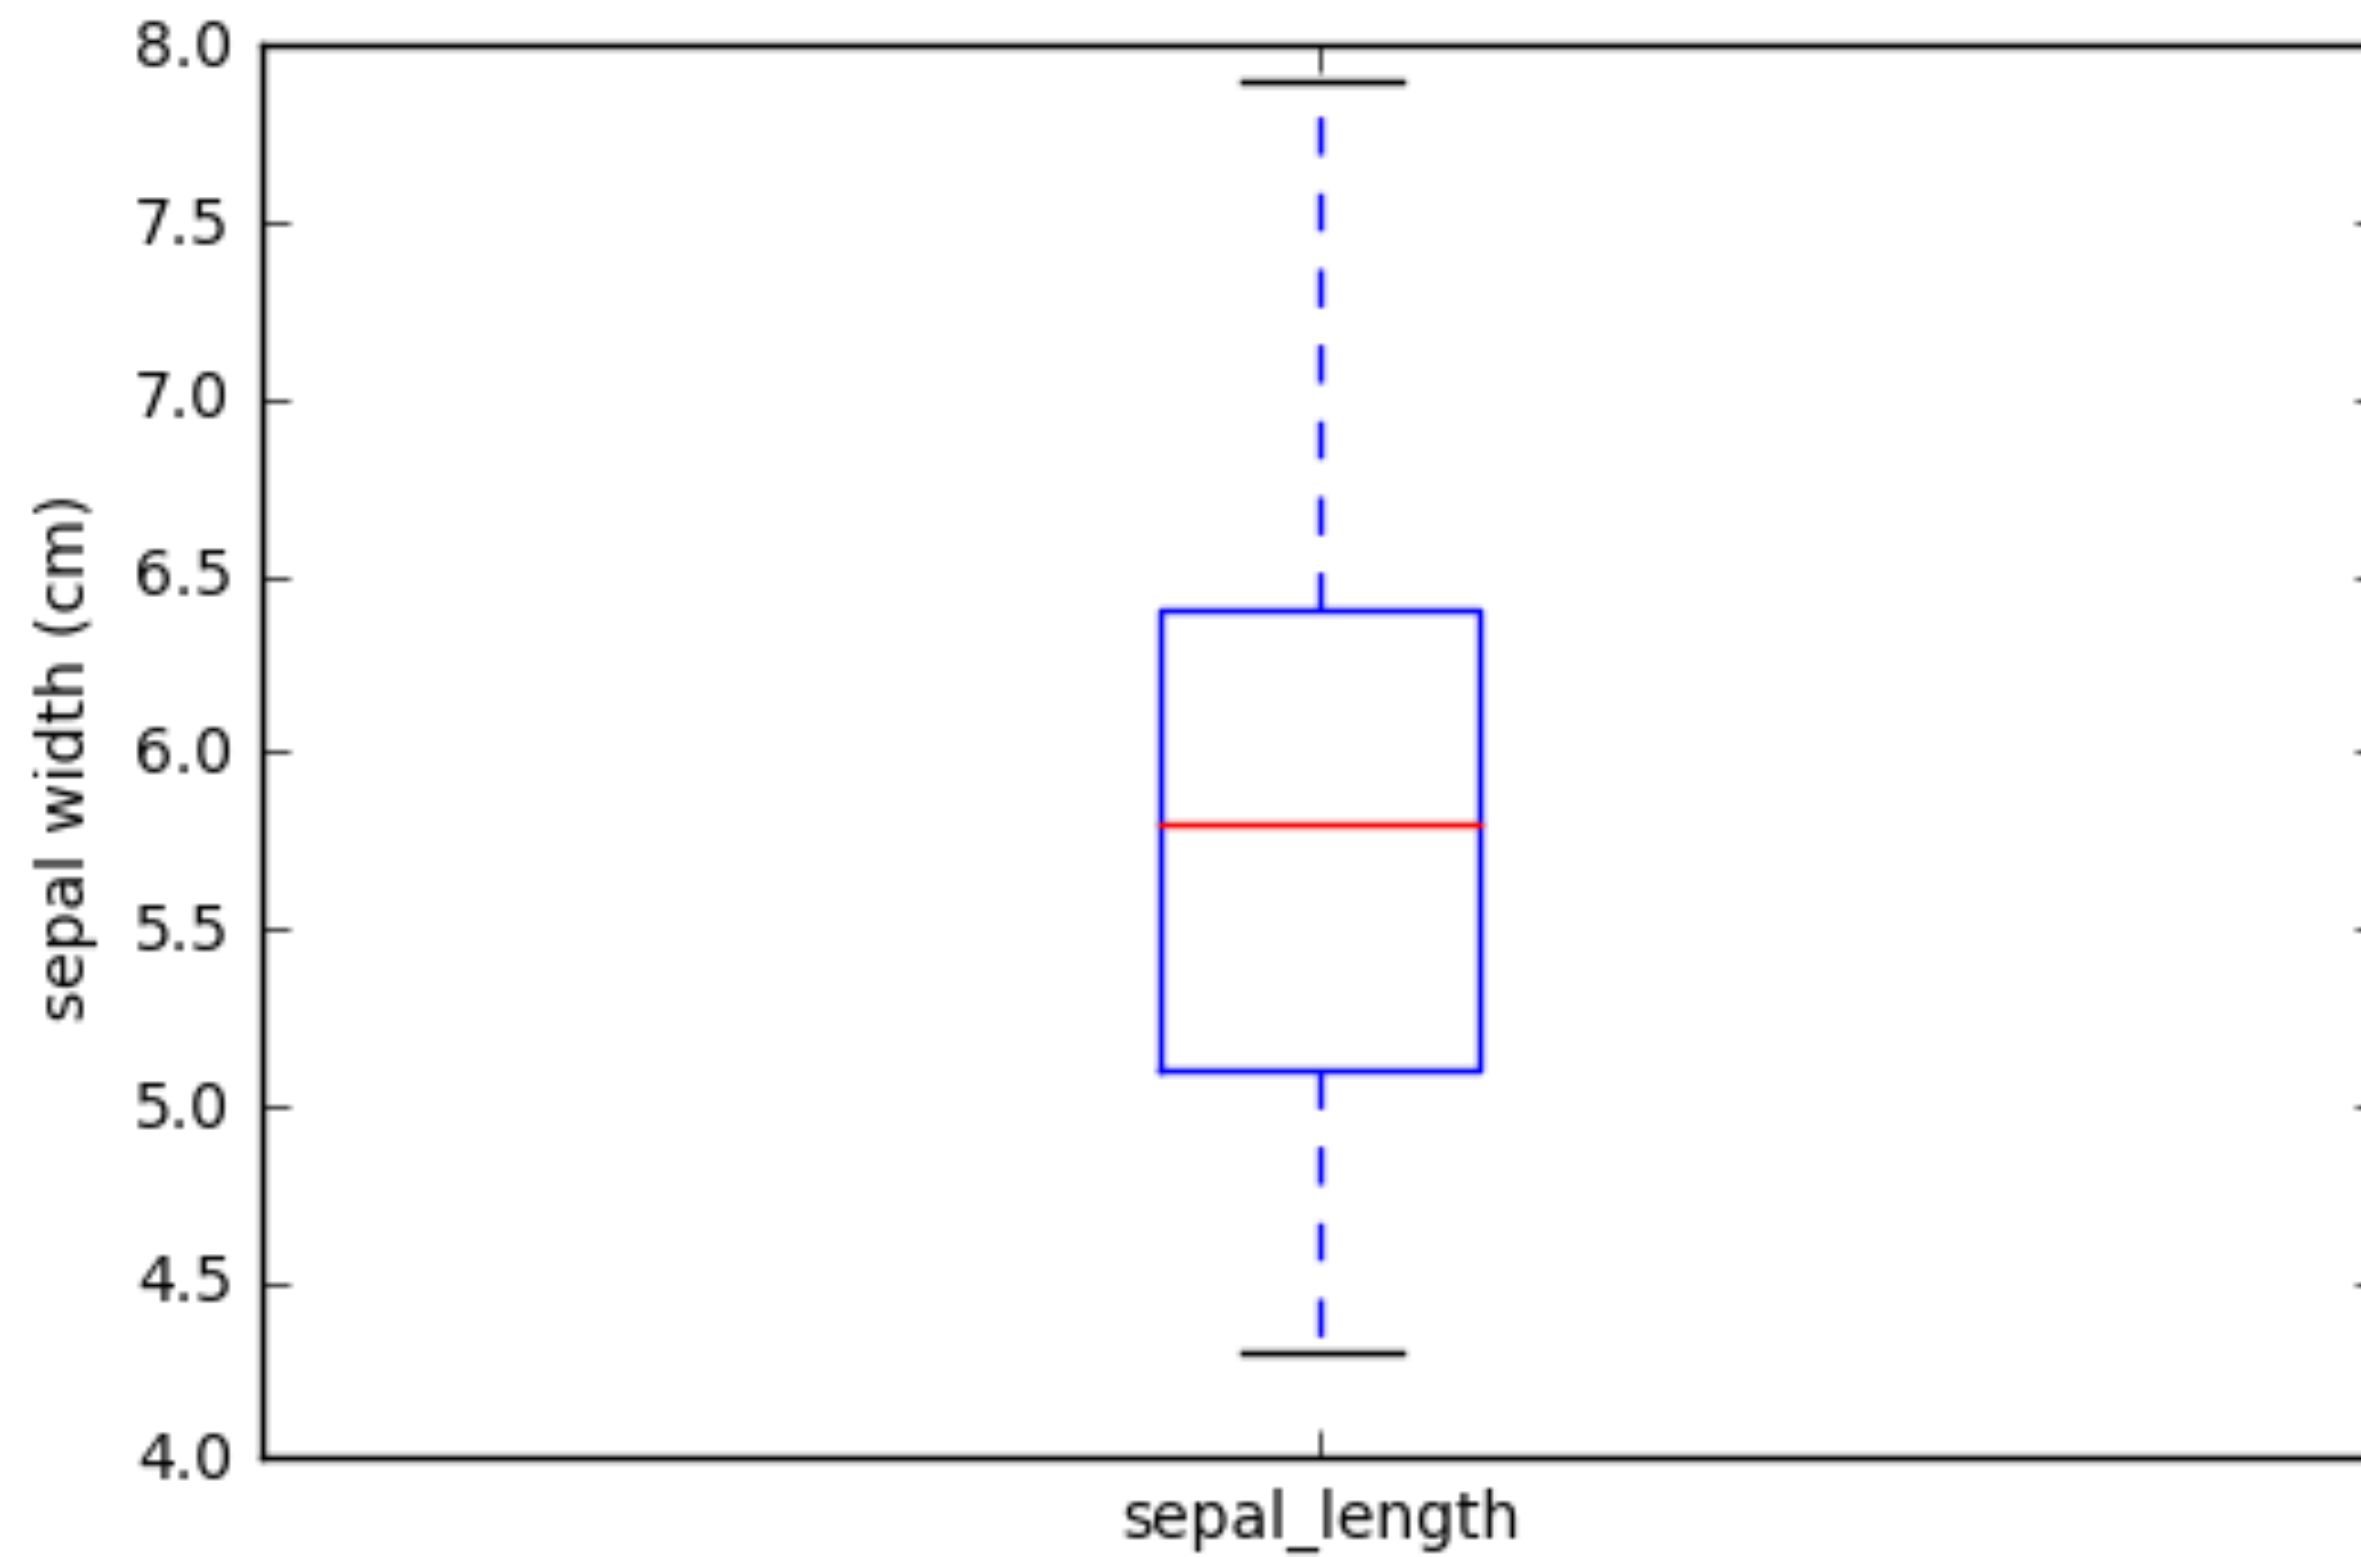
```
In [12]: iris.plot(y='sepal_length', kind='box')

In [13]: plt.ylabel('sepal width (cm)')

In [14]: plt.show()
```

Box plot

Histogram
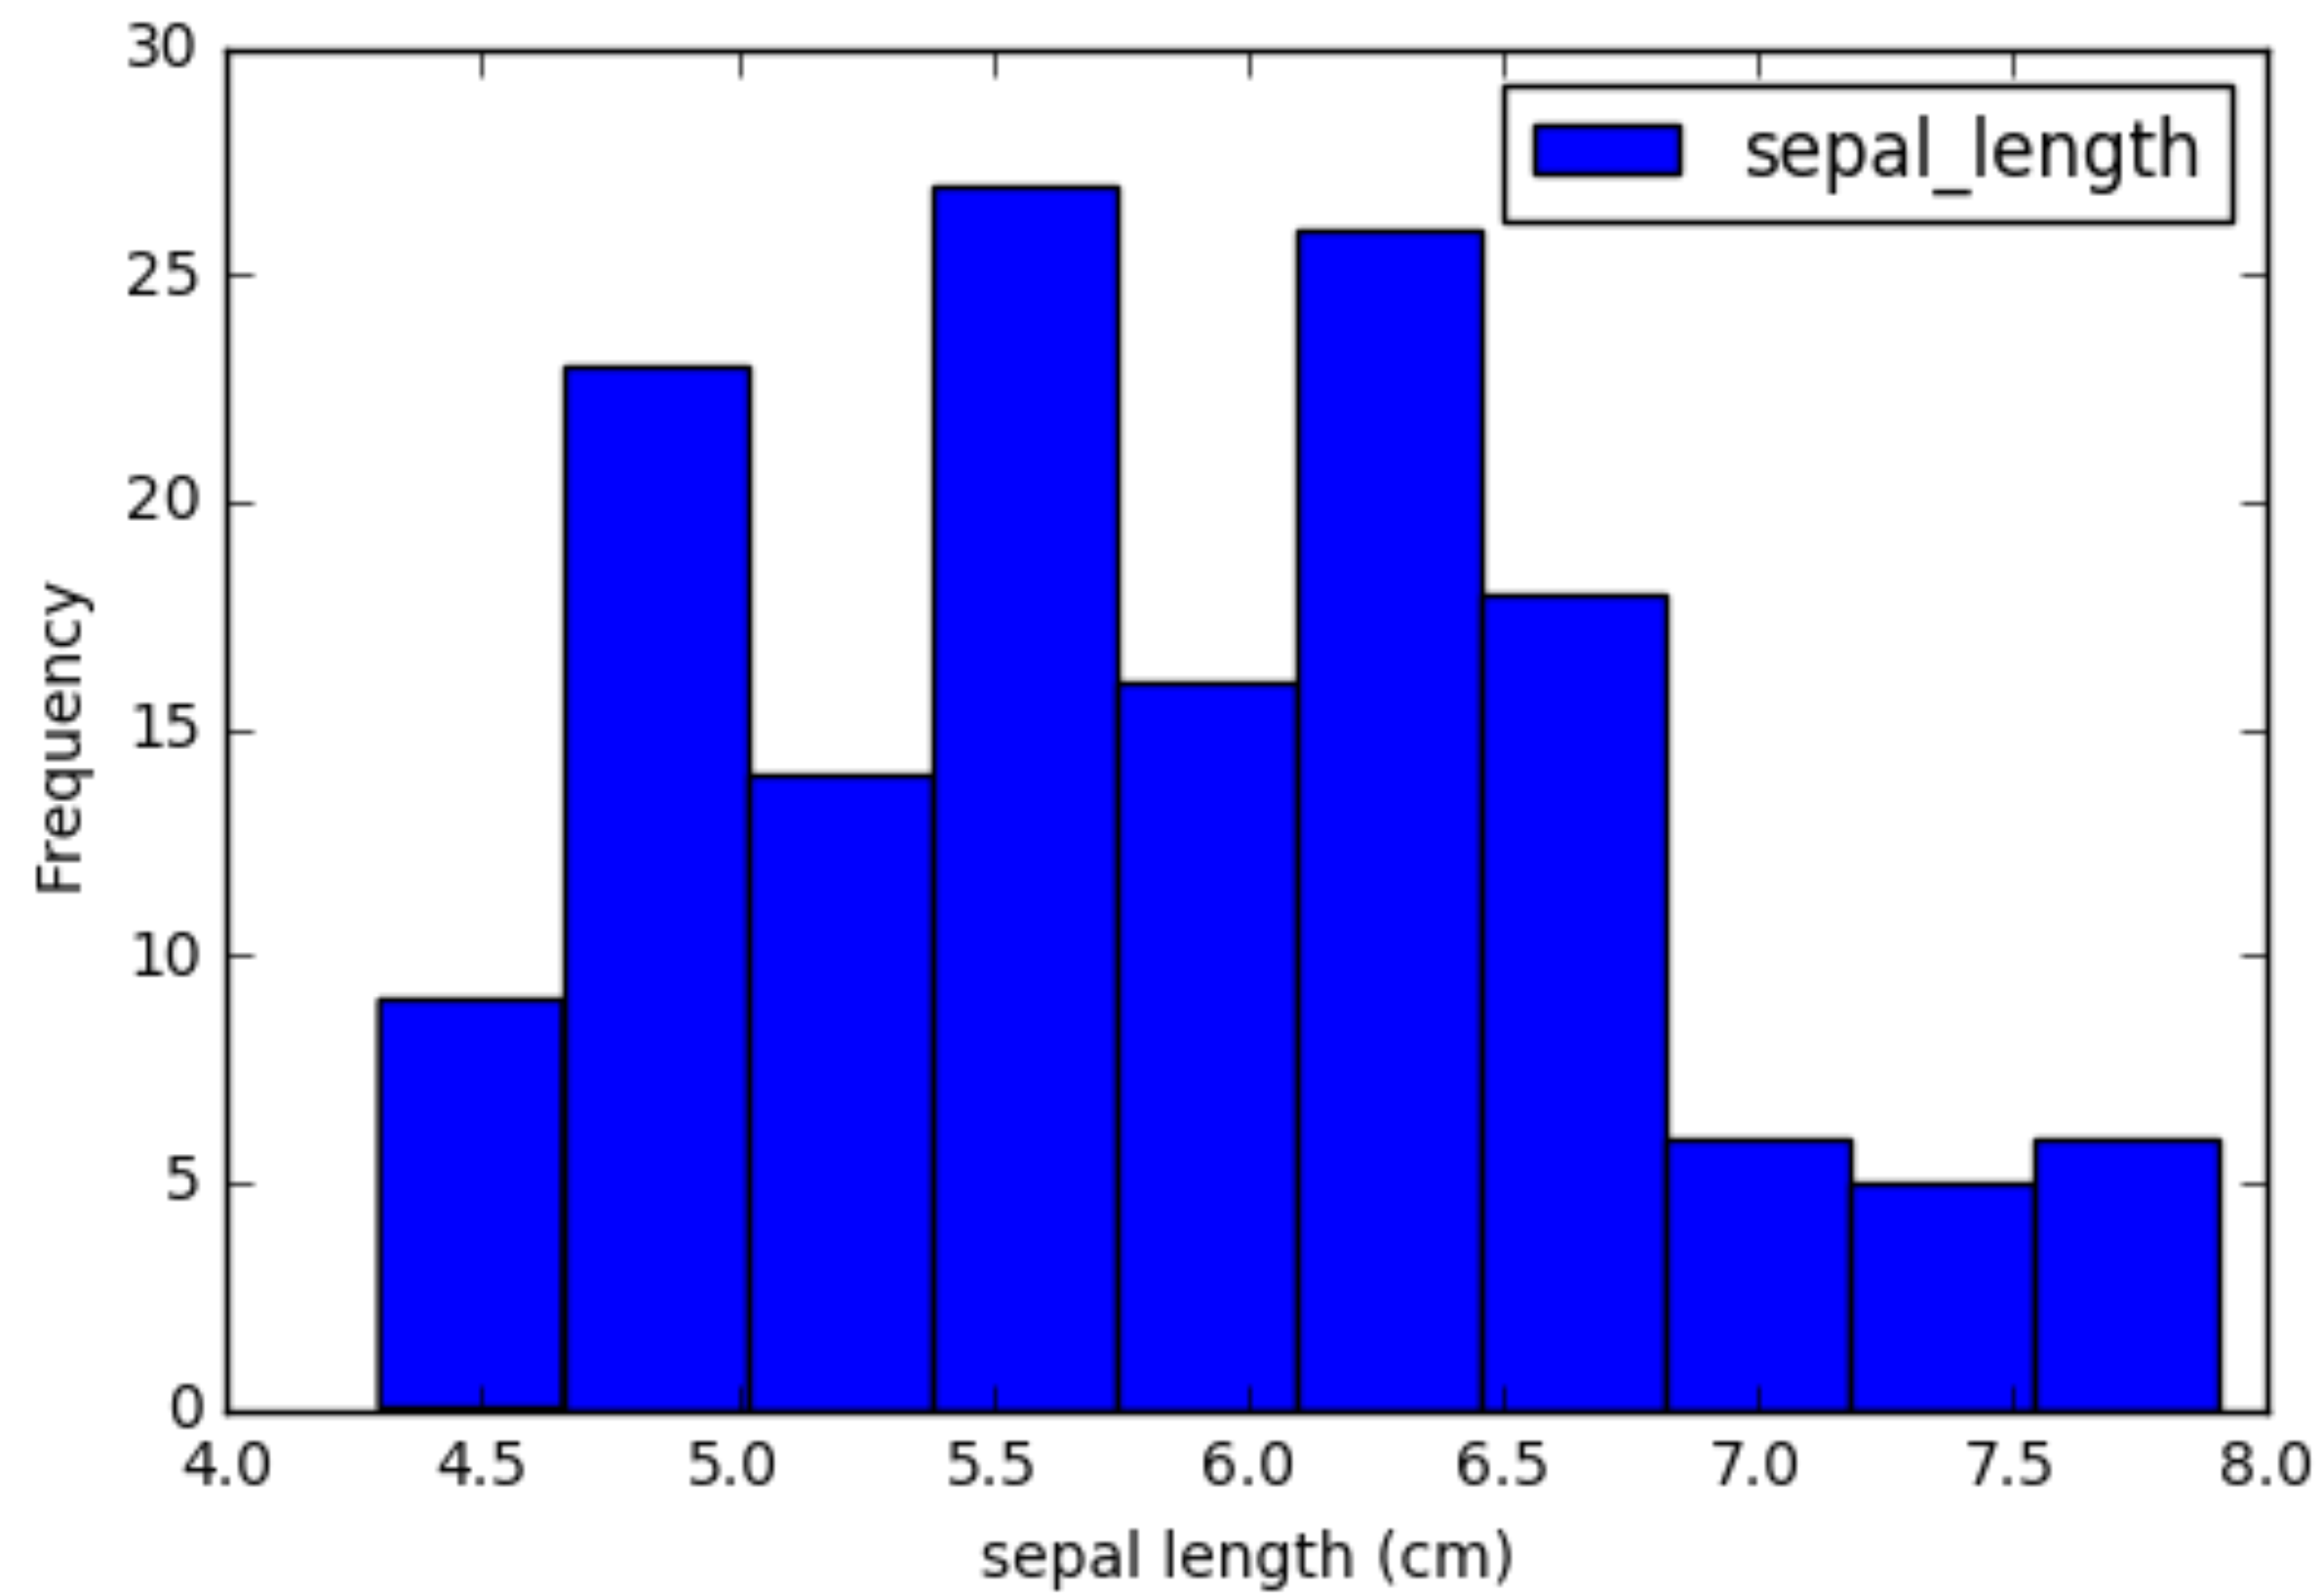
```
In [15]: iris.plot(y='sepal_length', kind='hist')

In [16]: plt.xlabel('sepal length (cm)')

In [17]: plt.show()
```

Histogram

Histogram options

- *bins* (integer): number of intervals or bins

- *range* (tuple): extrema of bins (minimum, maximum)

- *normed* (boolean): whether to normalize to one

- *cumulative* (boolean): compute Cumulative Distribution Function (CDF)

- ... more Matplotlib customizations
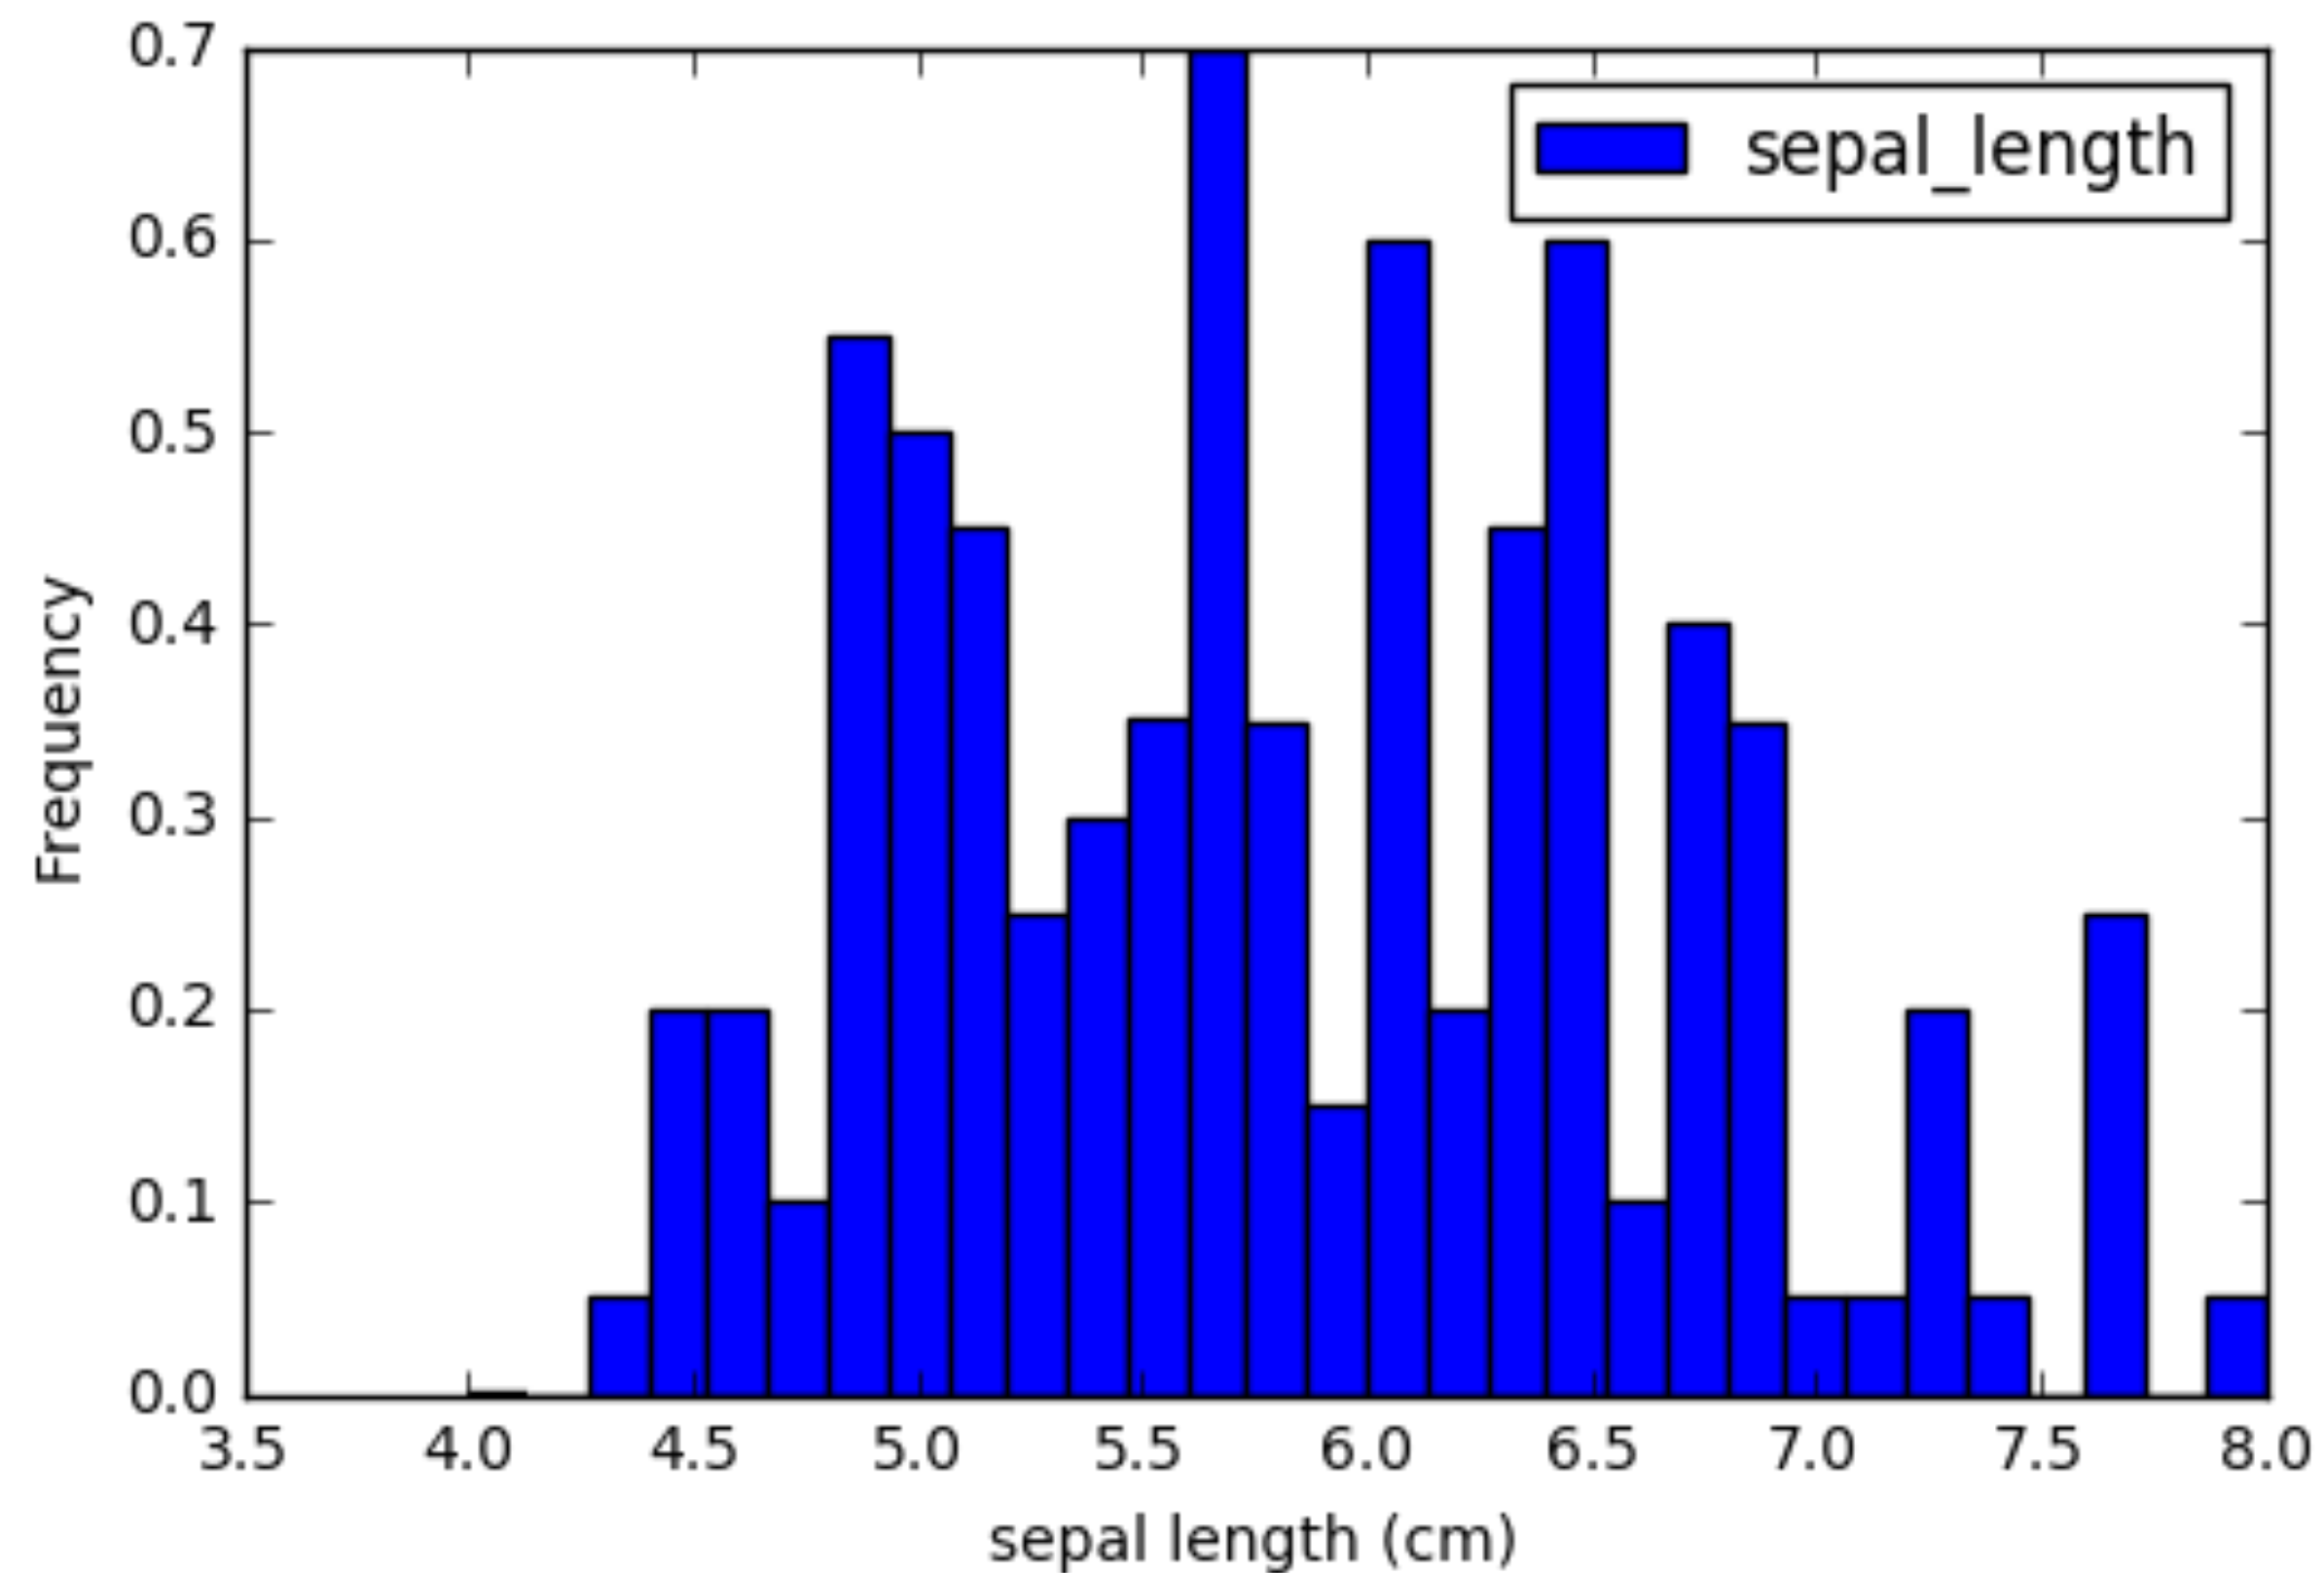
Customizing histogram

```
In [18]: iris.plot(y='sepal_length', kind='hist',
    ...:                  bins=30, range=(4,8), normed=True)

In [19]: plt.xlabel('sepal length (cm)')

In [20]: plt.show()
```

Customizing histogram

Cumulative distribution

```
In [21]: iris.plot(y='sepal_length', kind='hist', bins=30,
    ...:                 range=(4,8), cumulative=True, normed=True)

In [22]: plt.xlabel('sepal length (cm)')

In [23]: plt.title('Cumulative distribution function (CDF)')

In [24]: plt.show()
```

Cumulative distribution
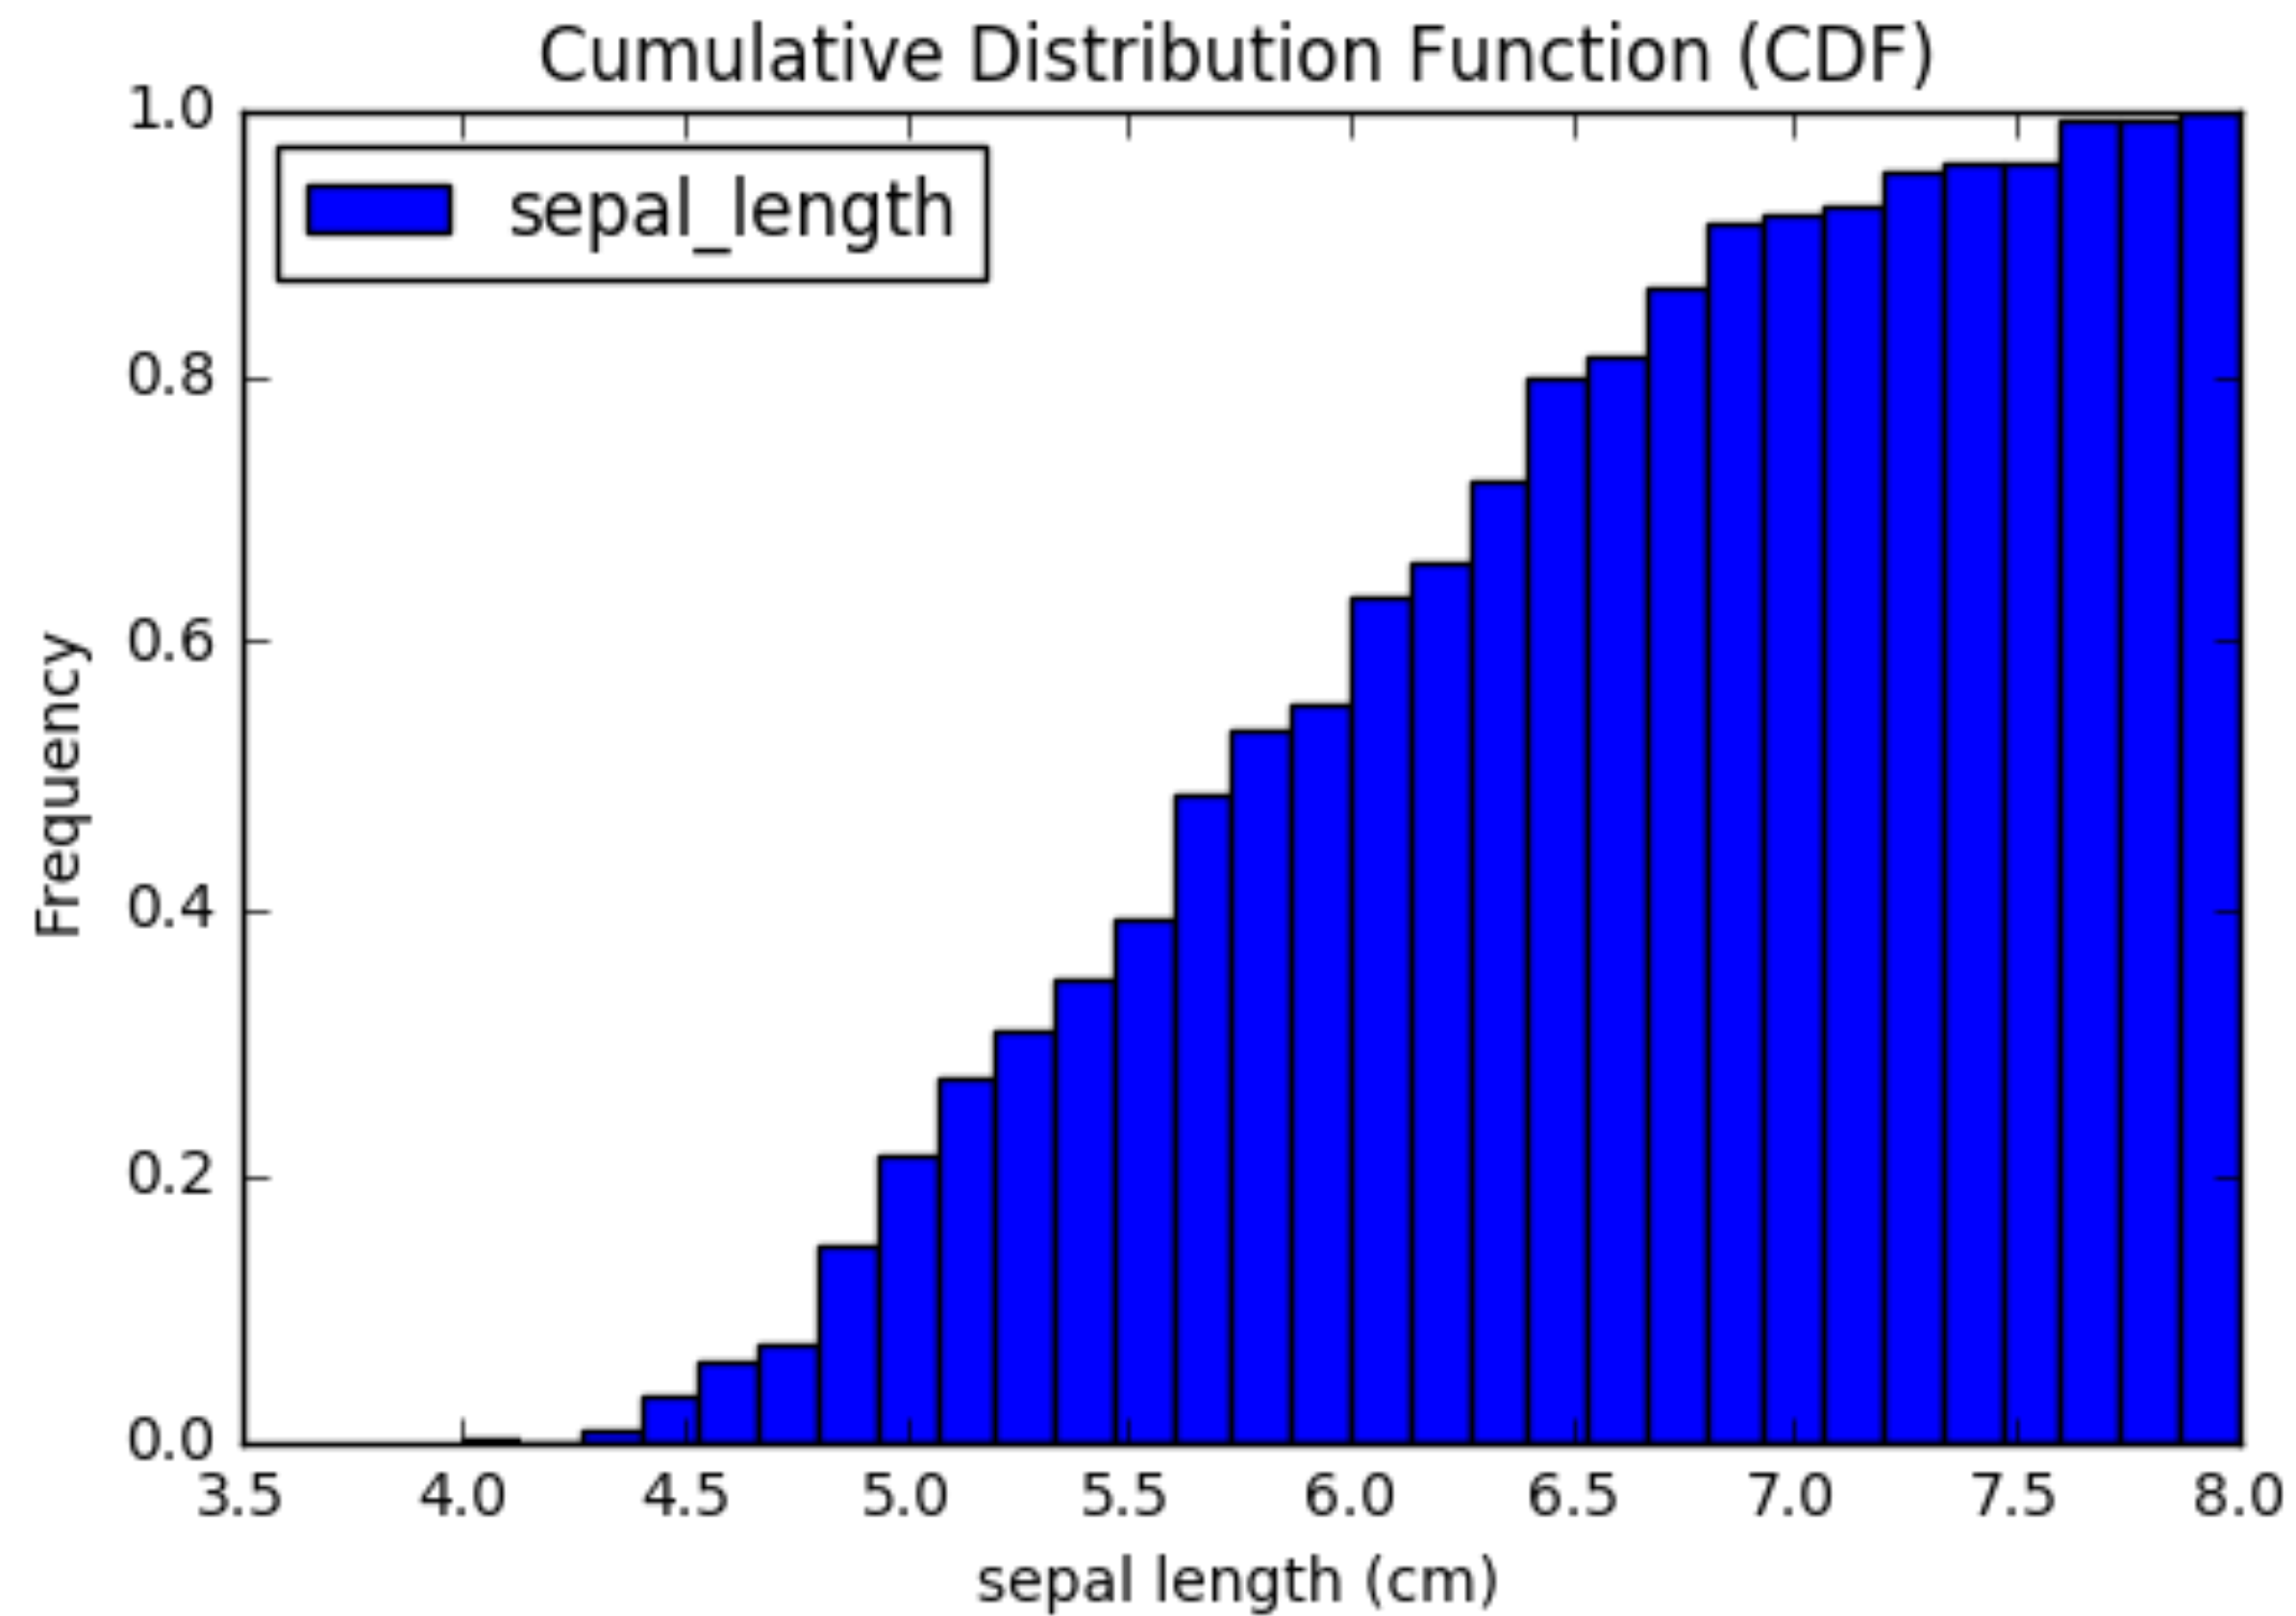
Word of warning

- Three different DataFrame plot idioms

  - *iris.plot(kind='hist')*

  - *iris.plt.hist()*

  - *iris.hist()*

- Syntax/results differ!

- Pandas API still evolving: check documentation!

PANDAS FOUNDATIONS

Let's practice!

Statistical exploratory data analysis

Summarizing with describe()

```
In [1]: iris.describe() # summary statistics
Out[1]:
        sepal_length   sepal_width   petal_length   petal_width
count    150.000000    150.000000    150.000000    150.000000
mean       5.843333      3.057333      3.758000      1.199333
std        0.828066      0.435866      1.765298      0.762238
min        4.300000      2.000000      1.000000      0.100000
25%        5.100000      2.800000      1.600000      0.300000
50%        5.800000      3.000000      4.350000      1.300000
75%        6.400000      3.300000      5.100000      1.800000
max        7.900000      4.400000      6.900000      2.500000
```

Describe

- *count*: number of entries

- *mean*: average of entries

- *std*: standard deviation

- min: minimum entry

- *25%*: first quartile

- *50%*: median or second quartile

- *75%*: third quartile

- *max*: maximum entry

Counts

```
In [2]: iris['sepal_length'].count() # Applied to Series
Out[2]: 150

In [3]: iris['sepal_width'].count() # Applied to Series
Out[3]: 150

In [4]: iris[['petal_length', 'petal_width']].count() # Applied
   ...: to DataFrame
Out[4]:
petal_length    150
petal_width     150
dtype: int64

In [5]: type(iris[['petal_length', 'petal_width']].count()) #
   ...: returns Series
Out[5]: pandas.core.series.Series
```

Averages

```
In [6]: iris['sepal_length'].mean() # Applied to Series
Out[6]: 5.843333333333335

In [7]: iris.mean() # Applied to entire DataFrame
Out[7]:
sepal_length    5.843333
sepal_width     3.057333
petal_length    3.758000
petal_width     1.199333
dtype: float64
```

Standard deviations

```
In [8]: iris.std()
Out[8]:
sepal_length    0.828066
sepal_width     0.435866
petal_length    1.765298
petal_width     0.762238
dtype: float64
```

Mean and standard deviation on a bell curve

Medians

```
In [9]: iris.median()
Out[9]:
sepal_length    5.80
sepal_width     3.00
petal_length    4.35
petal_width     1.30
dtype: float64
```

Medians & 0.5 quantiles

```
In [10]: iris.median()
Out[10]:
sepal_length    5.80
sepal_width     3.00
petal_length    4.35
petal_width     1.30
dtype: float64

In [11]: q = 0.5

In [12]: iris.quantile(q)
Out[12]:
sepal_length    5.80
sepal_width     3.00
petal_length    4.35
petal_width     1.30
dtype: float64
```

Inter-quartile range (IQR)

```
In [13]: q = [0.25, 0.75]

In [14]: iris.quantile(q)
Out[14]:
      sepal_length  sepal_width  petal_length  petal_width
0.25           5.1          2.8           1.6          0.3
0.75           6.4          3.3           5.1          1.8
```

Ranges

```
In [15]: iris.min()
Out[15]:
sepal_length        4.3
sepal_width           2
petal_length          1
petal_width         0.1
species          setosa
dtype: object

In [16]: iris.max()
Out[16]:
sepal_length          7.9
sepal_width           4.4
petal_length          6.9
petal_width           2.5
species          virginica
dtype: object
```

Box plots

```
In [17]: iris.plot(kind= 'box')
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x118a3d5f8>

In [18]: plt.ylabel('[cm]')
Out[18]: <matplotlib.text.Text at 0x118a524e0>

In [19]: plt.show()
```
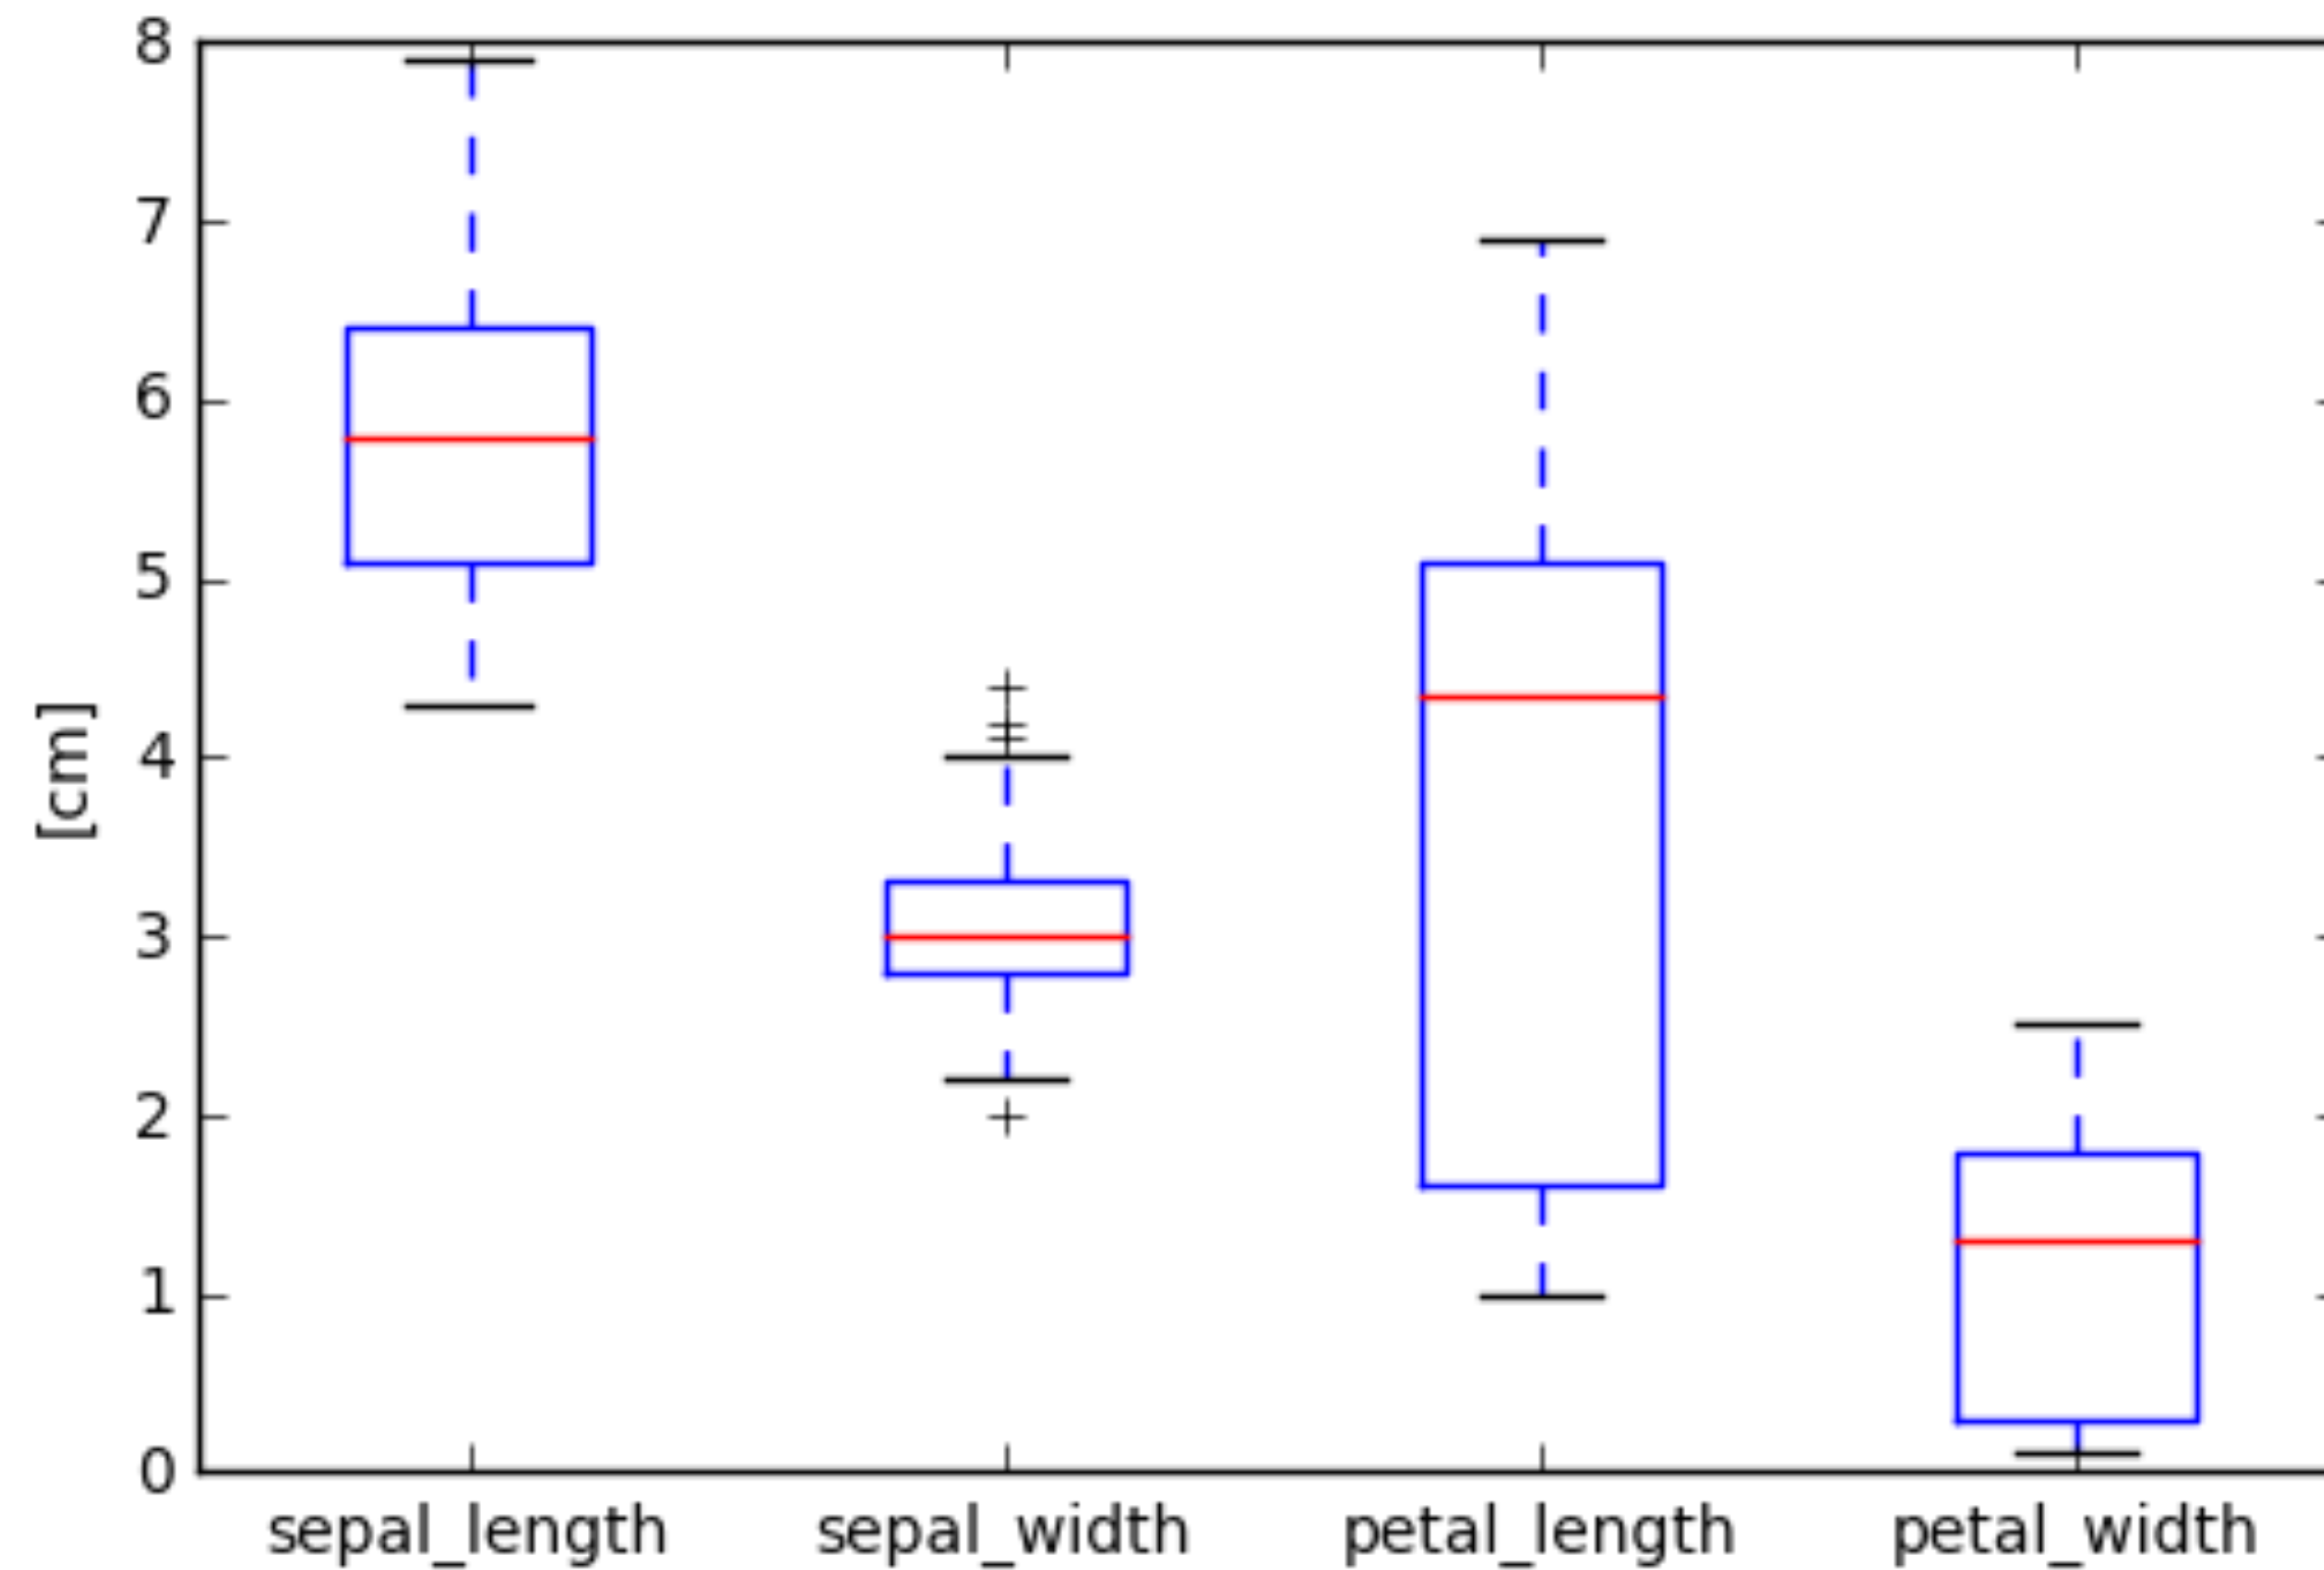
Box plots

Percentiles as quantiles

```
In [20]: iris.describe() # summary statistics
Out[20]:
        sepal_length    sepal_width    petal_length    petal_width
count     150.000000     150.000000      150.000000     150.000000
mean        5.843333       3.057333        3.758000       1.199333
std         0.828066       0.435866        1.765298       0.762238
min         4.300000       2.000000        1.000000       0.100000
25%         5.100000       2.800000        1.600000       0.300000
50%         5.800000       3.000000        4.350000       1.300000
75%         6.400000       3.300000        5.100000       1.800000
max         7.900000       4.400000        6.900000       2.500000
```

Let's practice!

Separating populations

```
In [1]: iris.head()
Out[1]:
   sepal_length  sepal_width  petal_length  petal_width  species
0           5.1          3.5           1.4          0.2   setosa
1           4.9          3.0           1.4          0.2   setosa
2           4.7          3.2           1.3          0.2   setosa
3           4.6          3.1           1.5          0.2   setosa
4           5.0          3.6           1.4          0.2   setosa
```

Describe species column

```
In [2]: iris['species'].describe()
Out[2]:
count             150
unique              3
top          setosa
freq               50
Name: species, dtype: object
```

count: # non-null entries

unique: # distinct values

top: most frequent category

freq: # occurrences of top

Unique & factors

```
In [3]: iris['species'].unique()
Out[3]: array(['setosa', 'versicolor', 'virginica'], dtype=object)
```

Filtering by species

```
In [4]: indices = iris['species'] == 'setosa'

In [5]: setosa = iris.loc[indices,:] # extract new DataFrame

In [6]: indices = iris['species'] == 'versicolor'

In [7]: versicolor = iris.loc[indices,:] # extract new DataFrame

In [8]: indices = iris['species'] == 'virginica'

In [9]: virginica = iris.loc[indices,:] # extract new DataFrame
```

Checking species

```
In [10]: setosa['species'].unique()
Out[10]: array(['setosa'], dtype=object)

In [11]: versicolor['species'].unique()
Out[11]: array(['versicolor'], dtype=object)

In [12]: virginica['species'].unique()
Out[12]: array(['virginica'], dtype=object)

In [13]: del setosa['species'], versicolor['species'],
    ...: virginica['species']
```

Checking indexes

```
In [14]: setosa.head(2)
Out[14]:
   sepal_length   sepal_width   petal_length   petal_width
0          5.1           3.5            1.4           0.2
1          4.9           3.0            1.4           0.2


In [15]: versicolor.head(2)
Out[15]:
    sepal_length   sepal_width   petal_length   petal_width
50           7.0           3.2            4.7           1.4
51           6.4           3.2            4.5           1.5


In [16]: virginica.head(2)
Out[16]:
     sepal_length   sepal_width   petal_length   petal_width
100           6.3           3.3            6.0           2.5
101           5.8           2.7            5.1           1.9
```

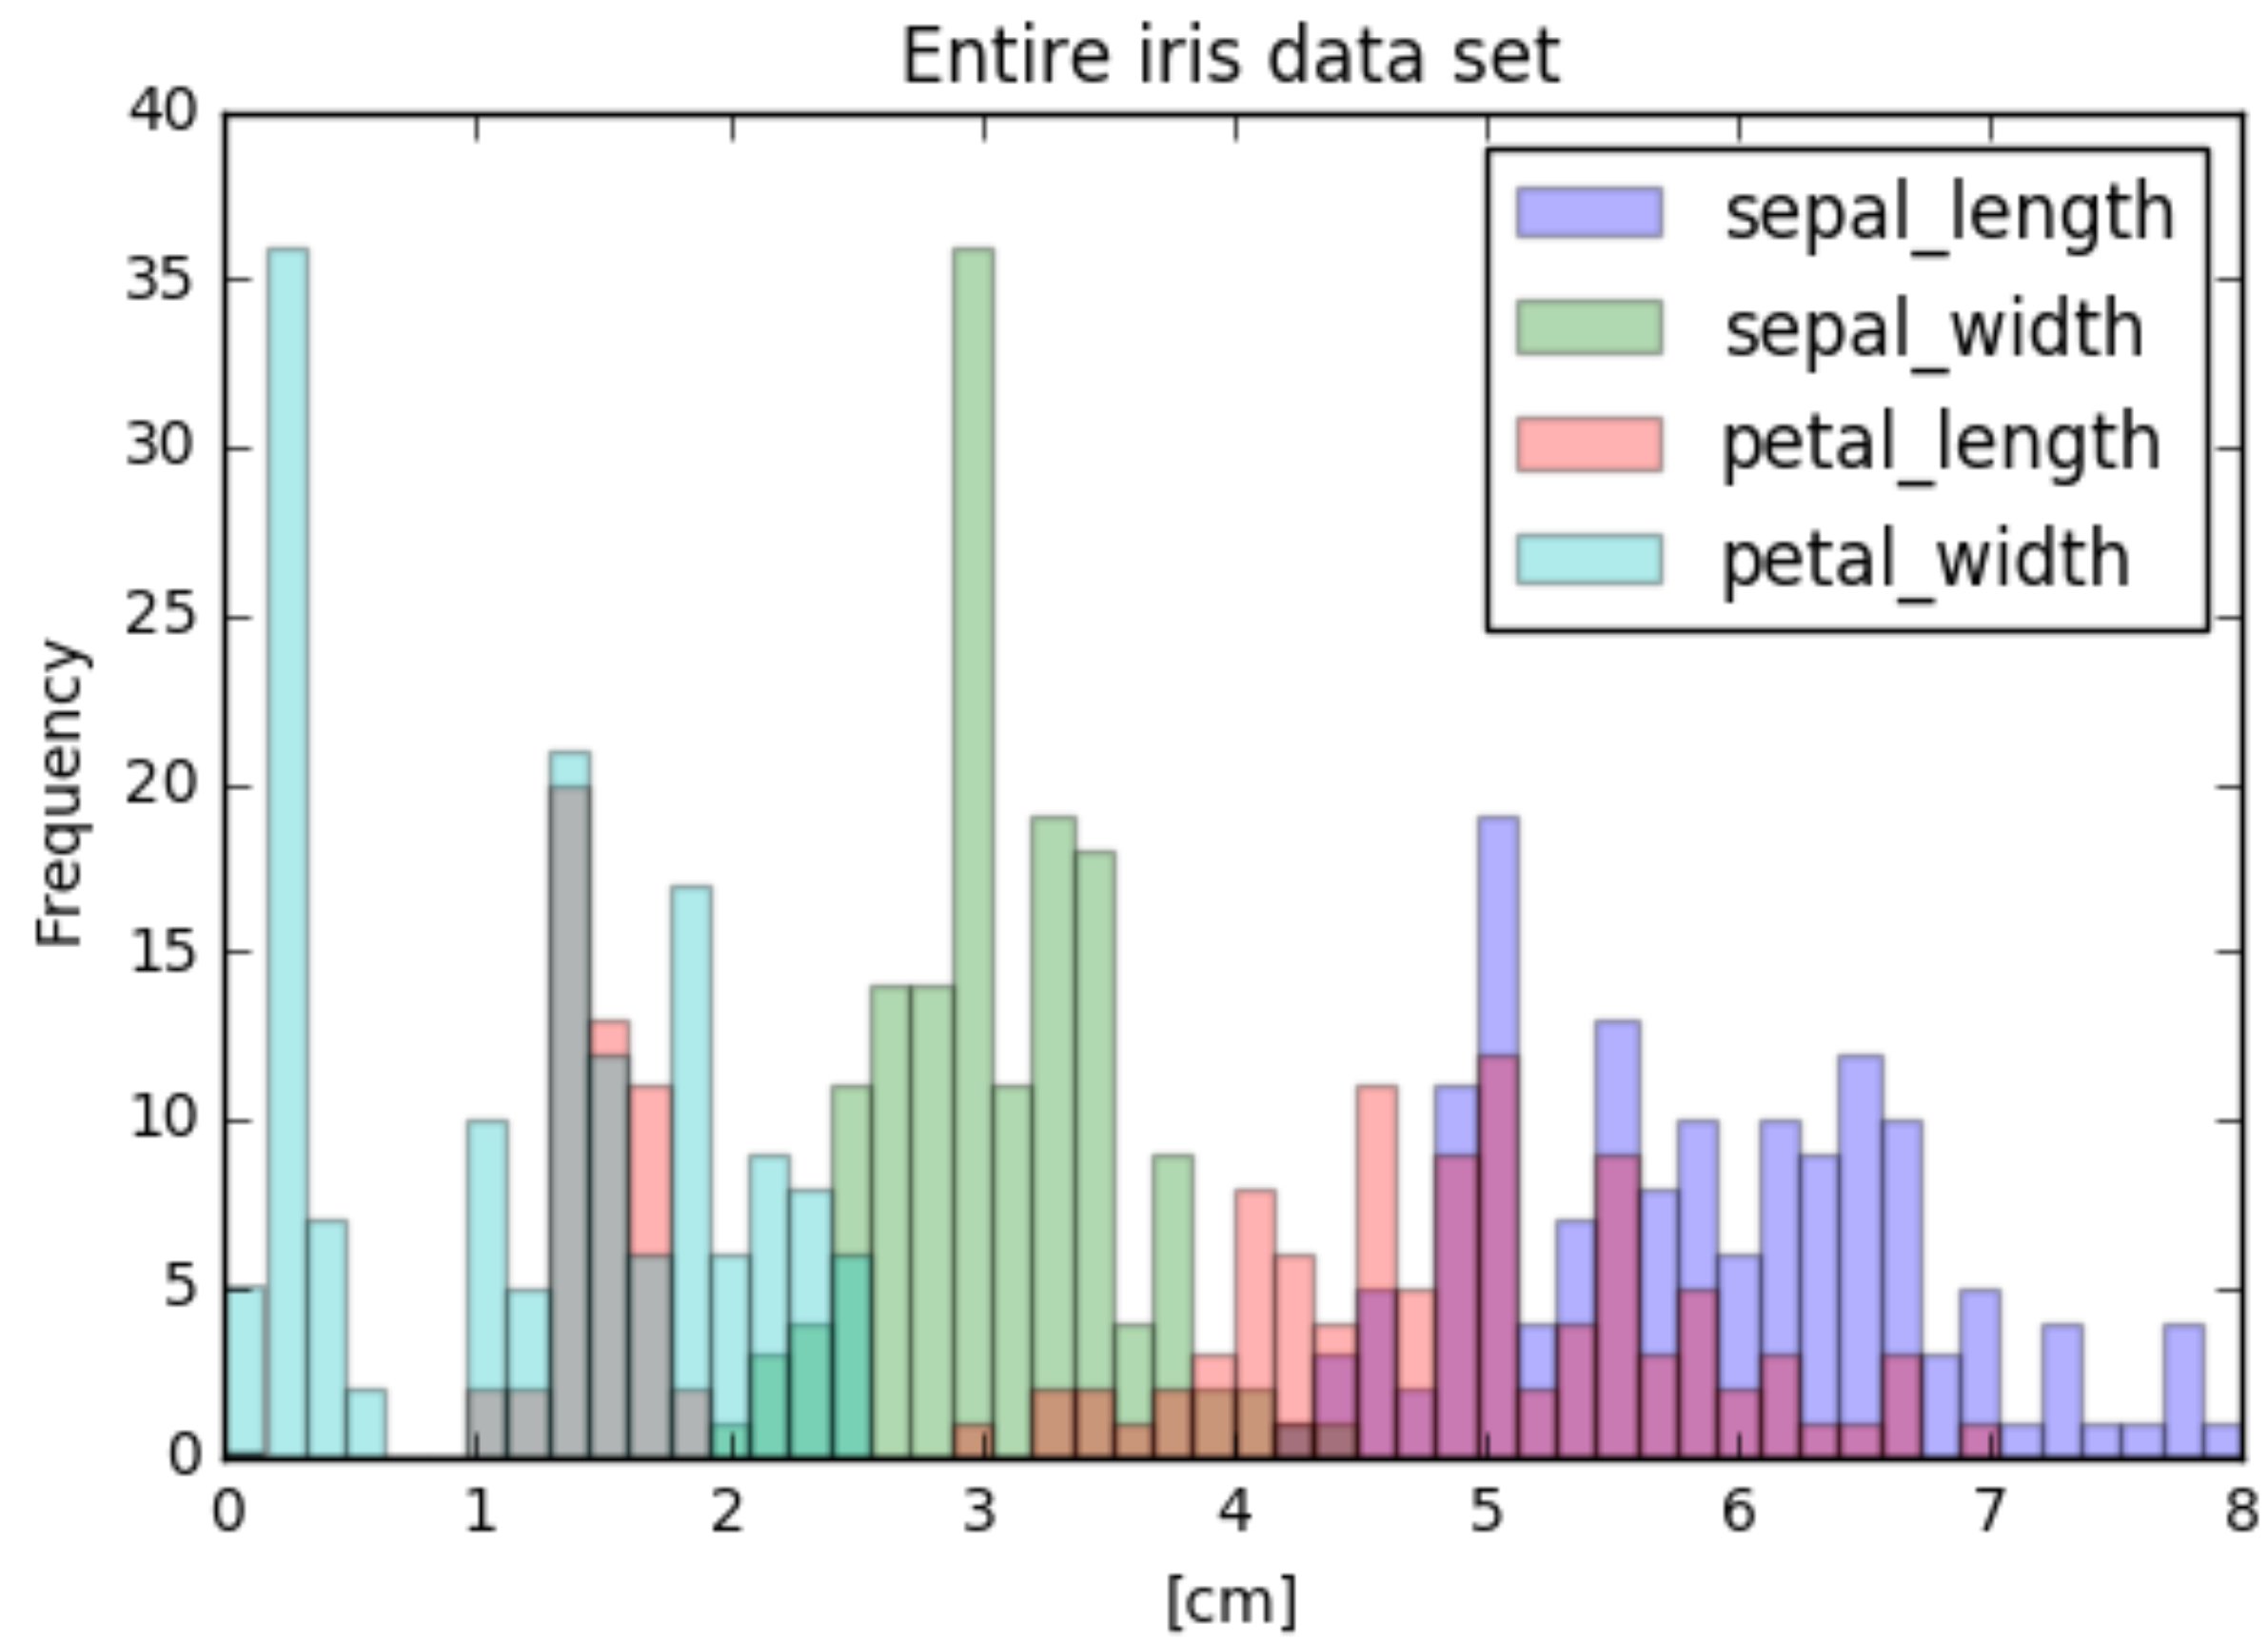# Visual EDA: all data

```
In [17]: iris.plot(kind= 'hist', bins=50, range=(0,8), alpha=0.3)

In [18]: plt.title('Entire iris data set')

In [19]: plt.xlabel('[cm]')

In [20]: plt.show()
```

Visual EDA: all data

Visual EDA: individual factors

```
In [21]: setosa.plot(kind='hist', bins=50, range=(0,8), alpha=0.3)

In [22]: plt.title('Setosa data set')

In [23]: plt.xlabel('[cm]')

In [24]: versicolor.plot(kind='hist', bins=50, range=(0,8), alpha=0.3)

In [25]: plt.title('Versicolor data set')

In [26]: plt.xlabel('[cm]')

In [27]: virginica.plot(kind='hist', bins=50, range=(0,8), alpha=0.3)

In [28]: plt.title('Virginica data set')

In [29]: plt.xlabel('[cm]')

In [30]: plt.show()
```
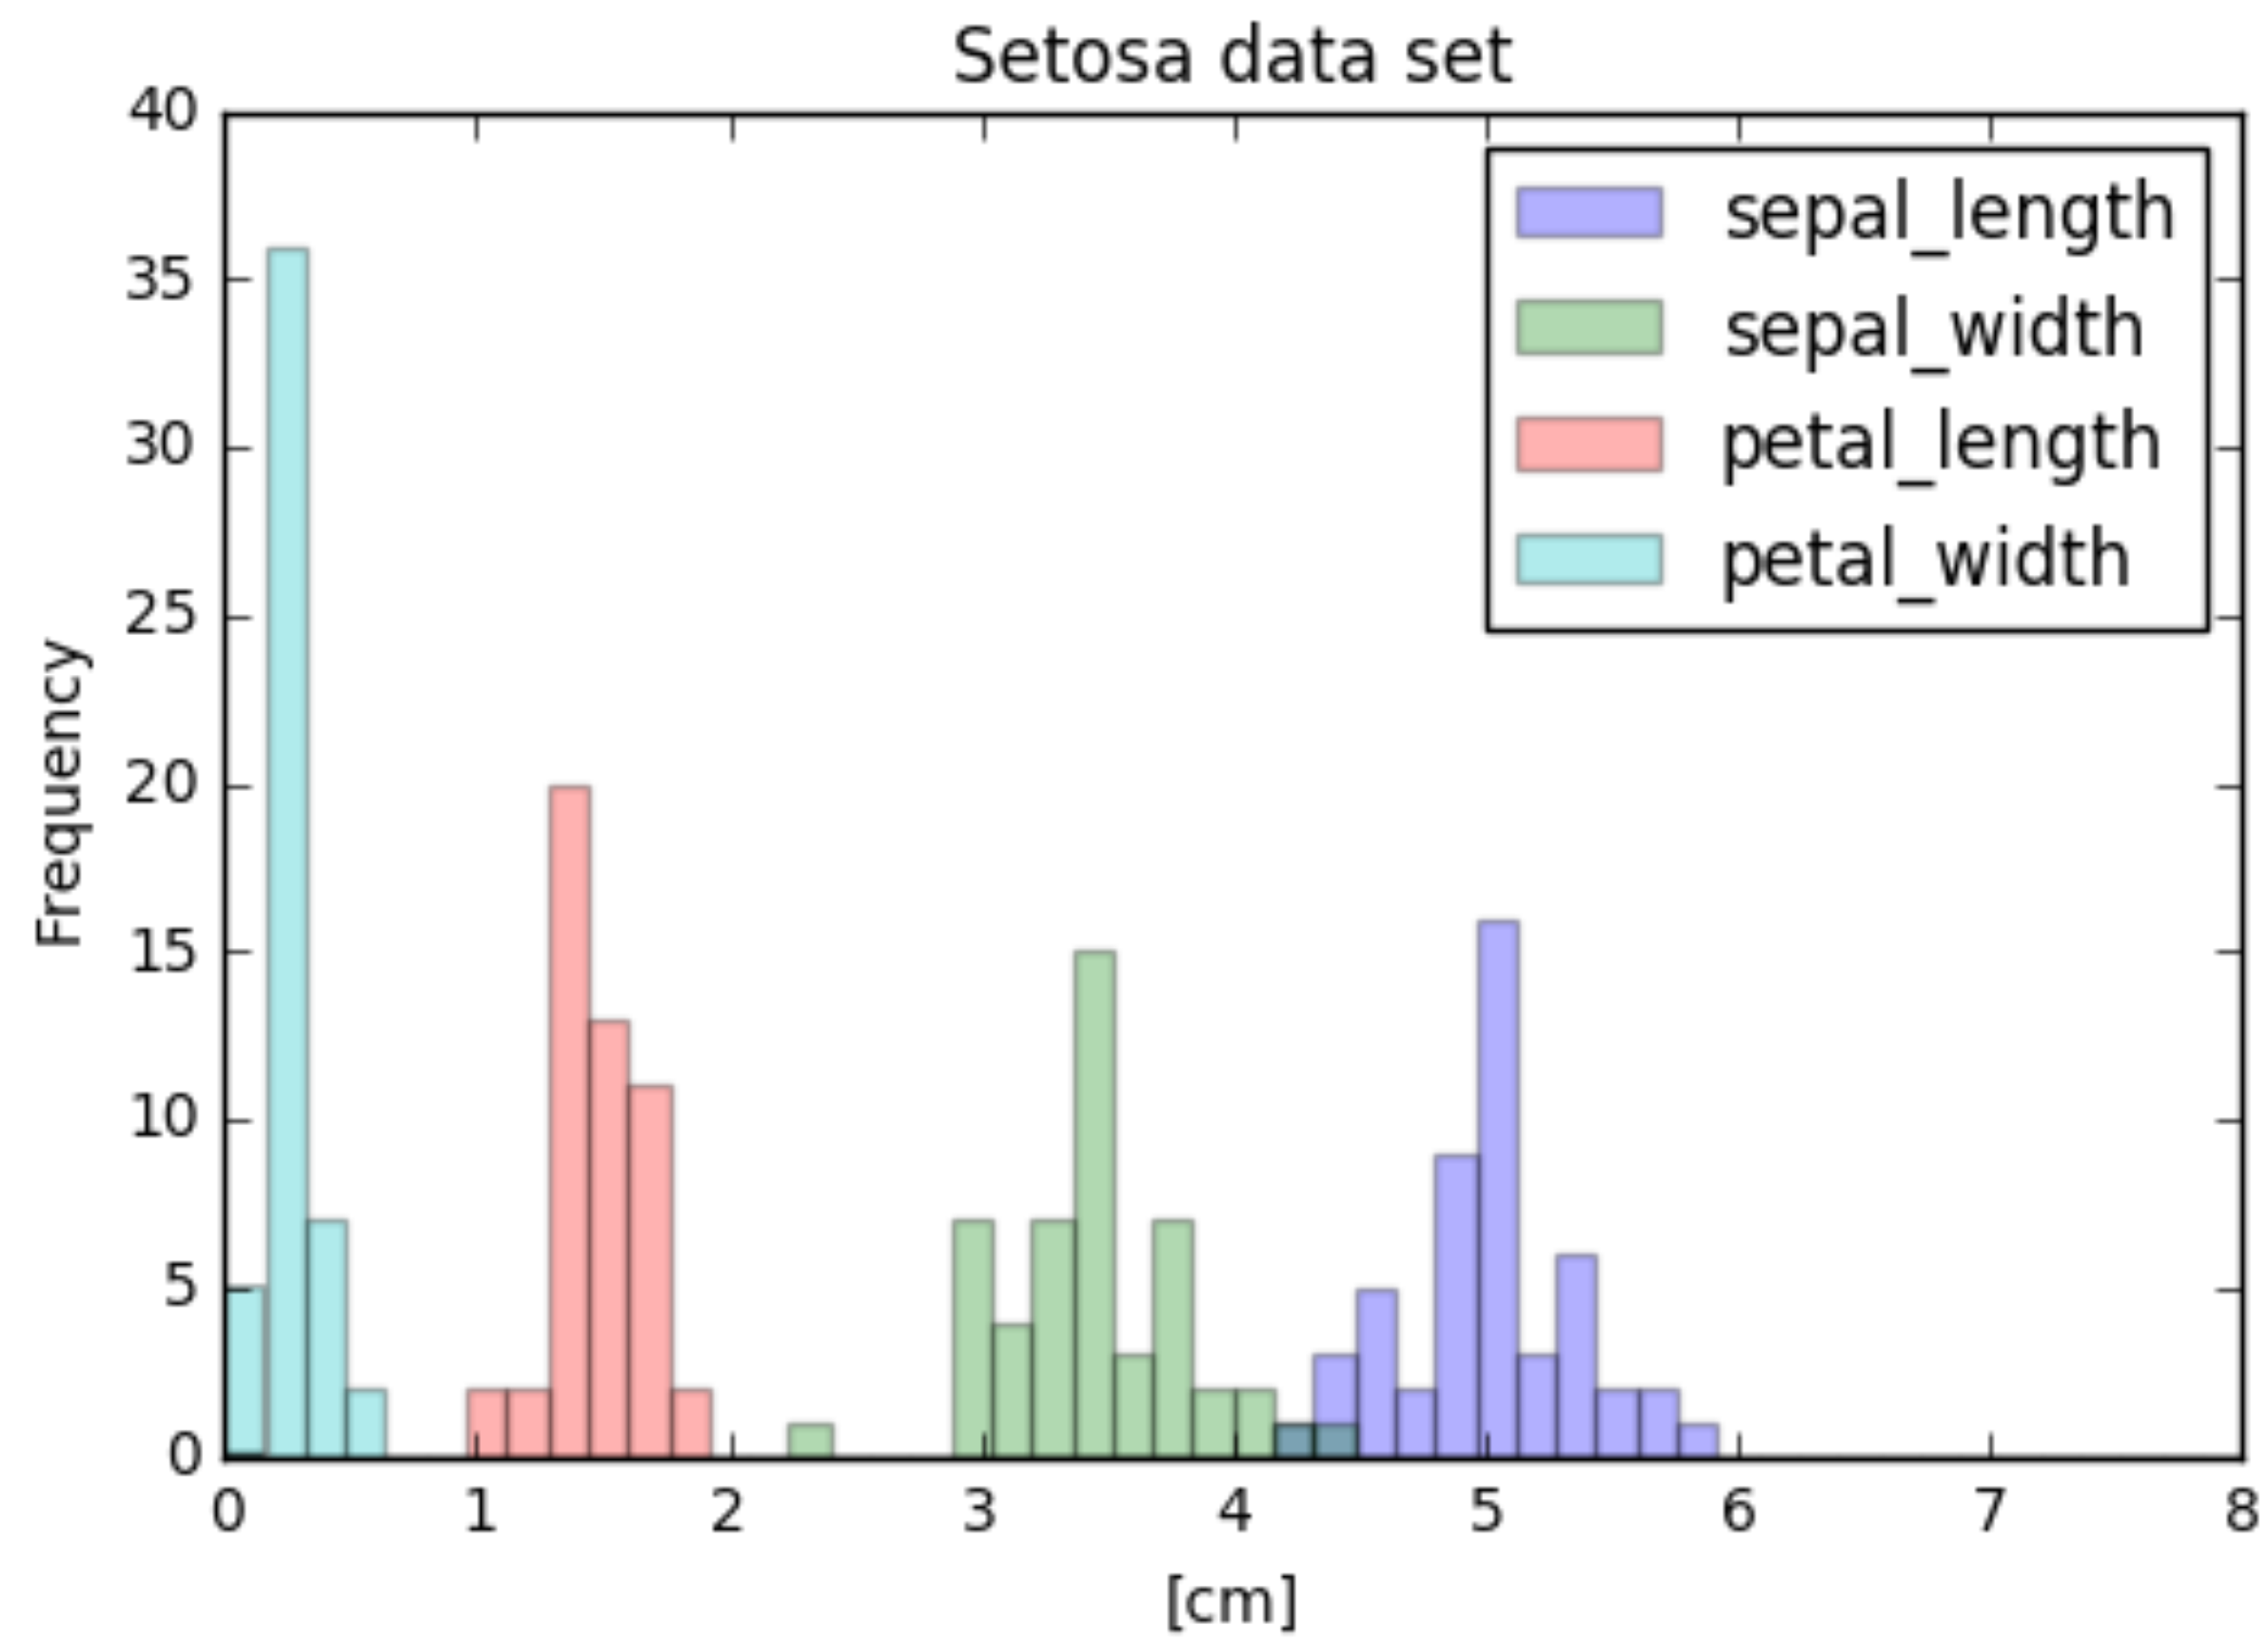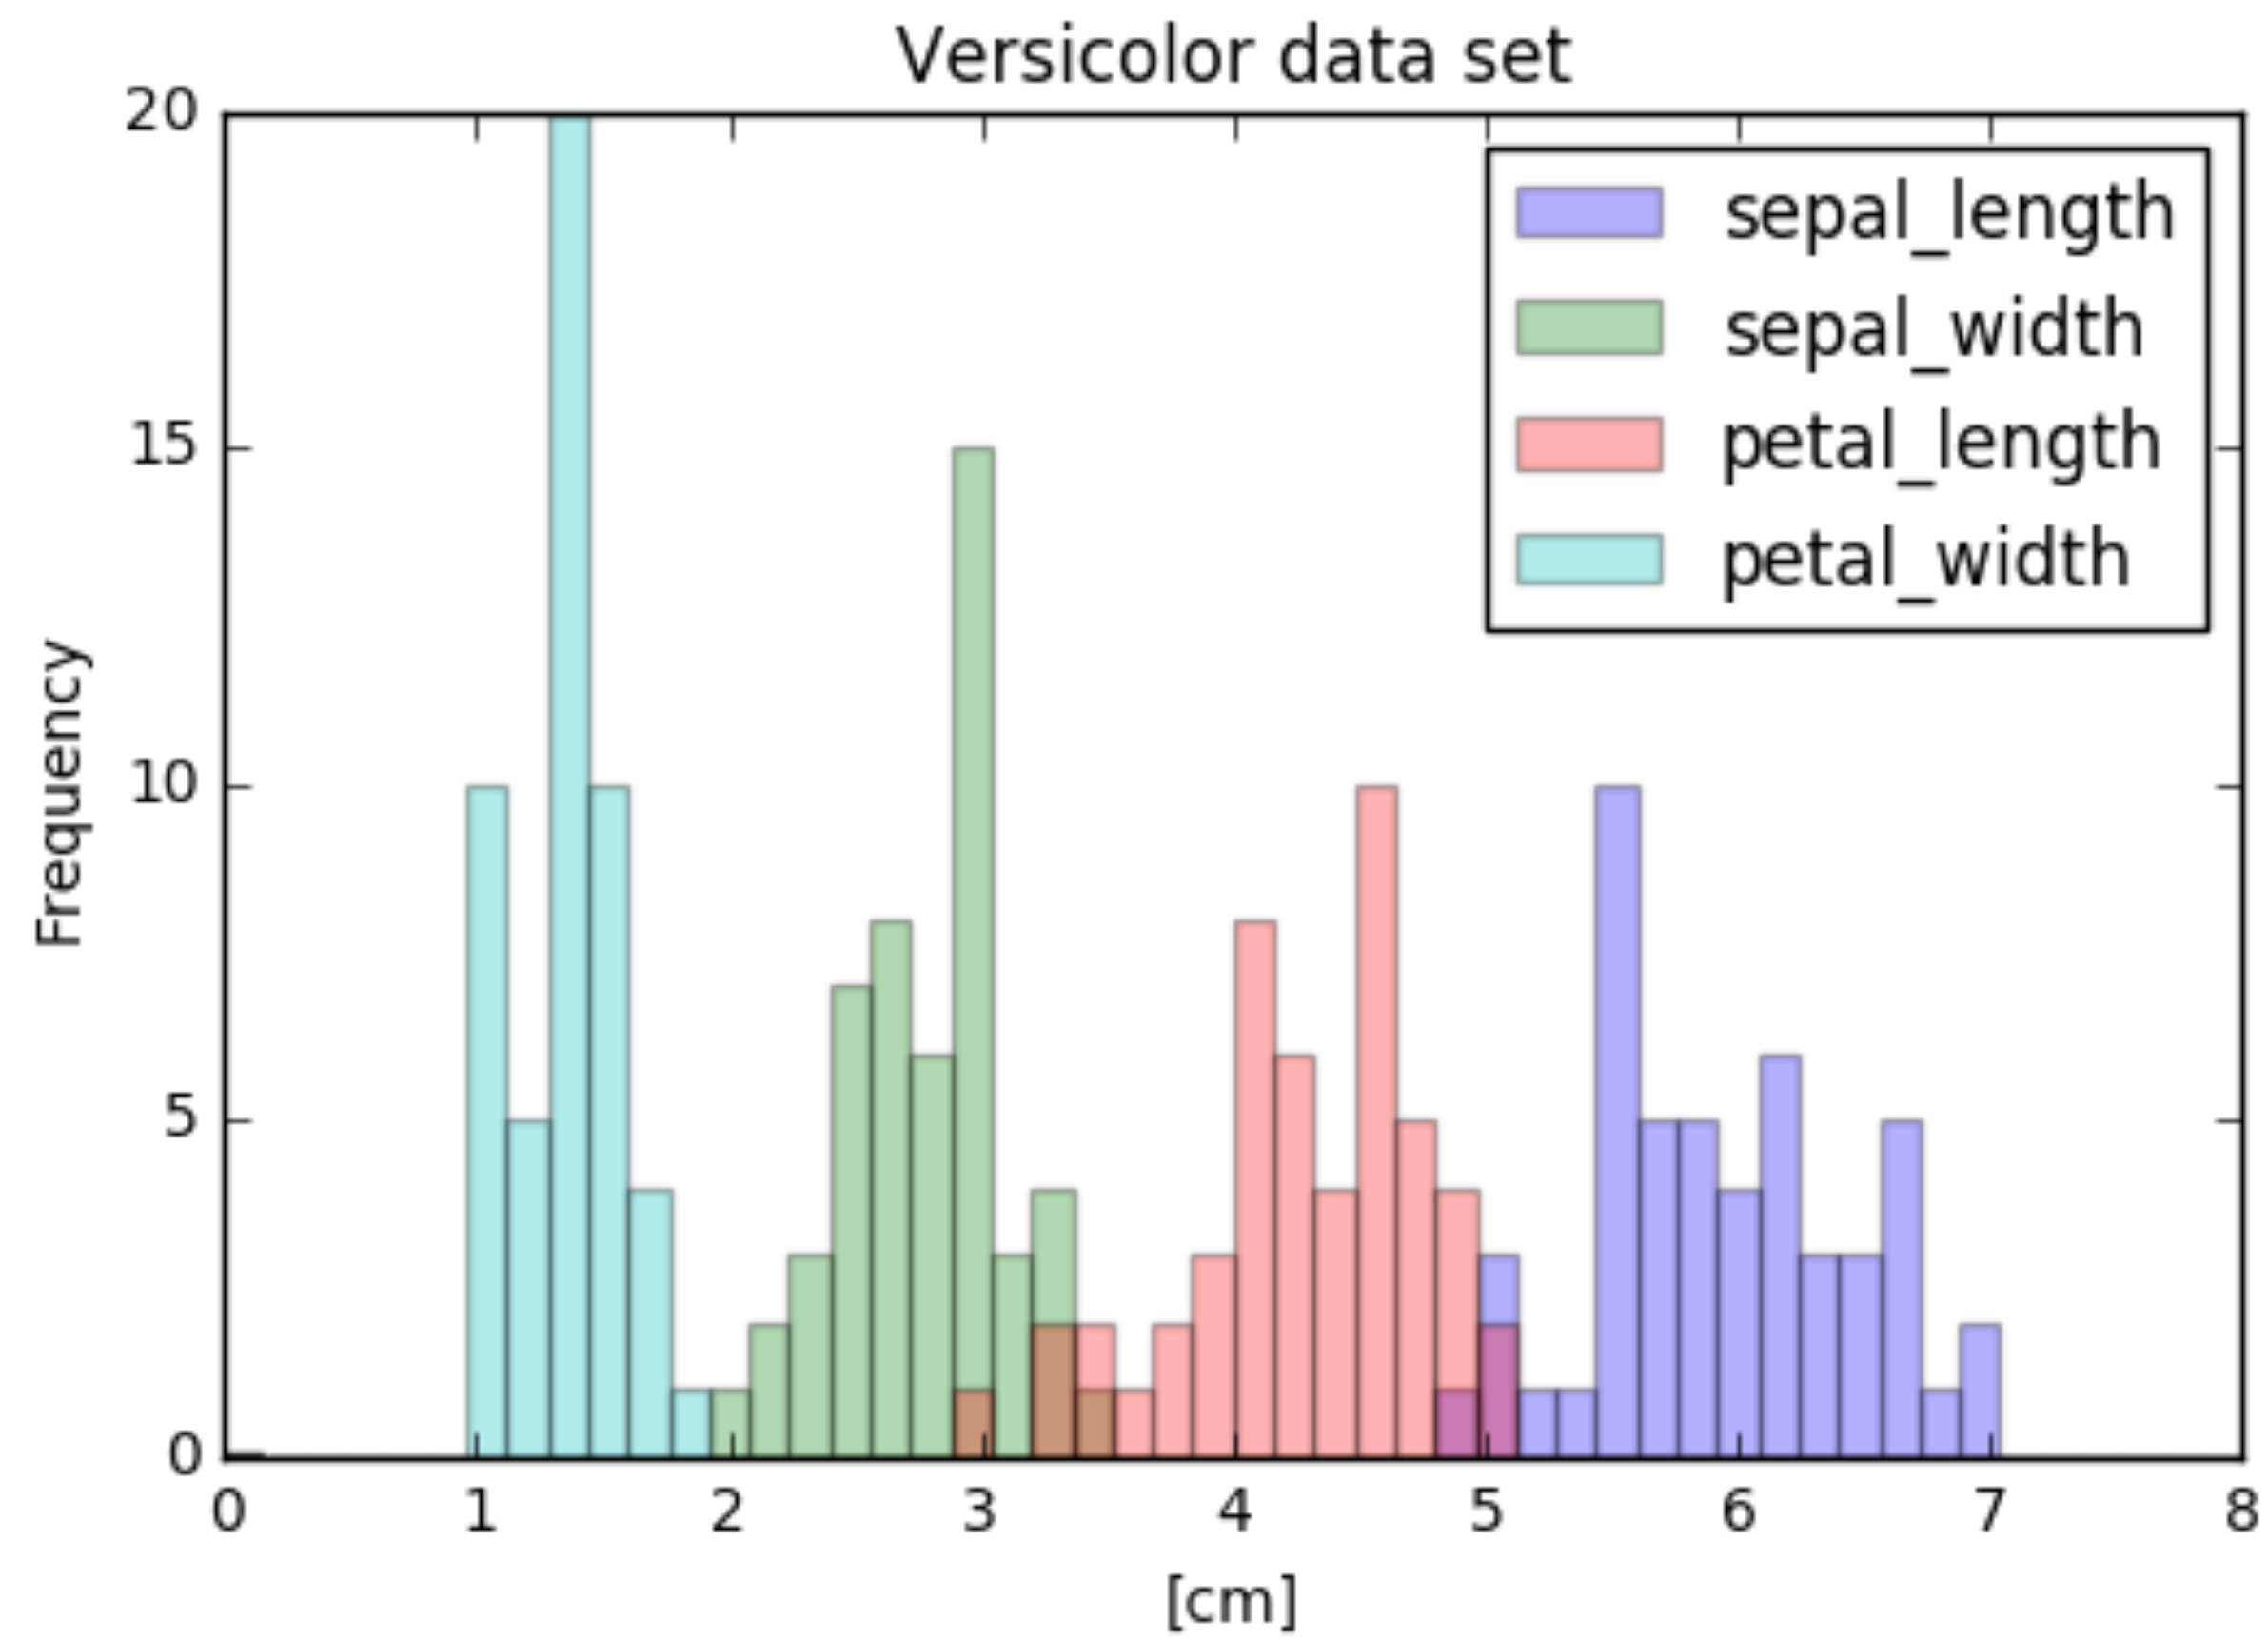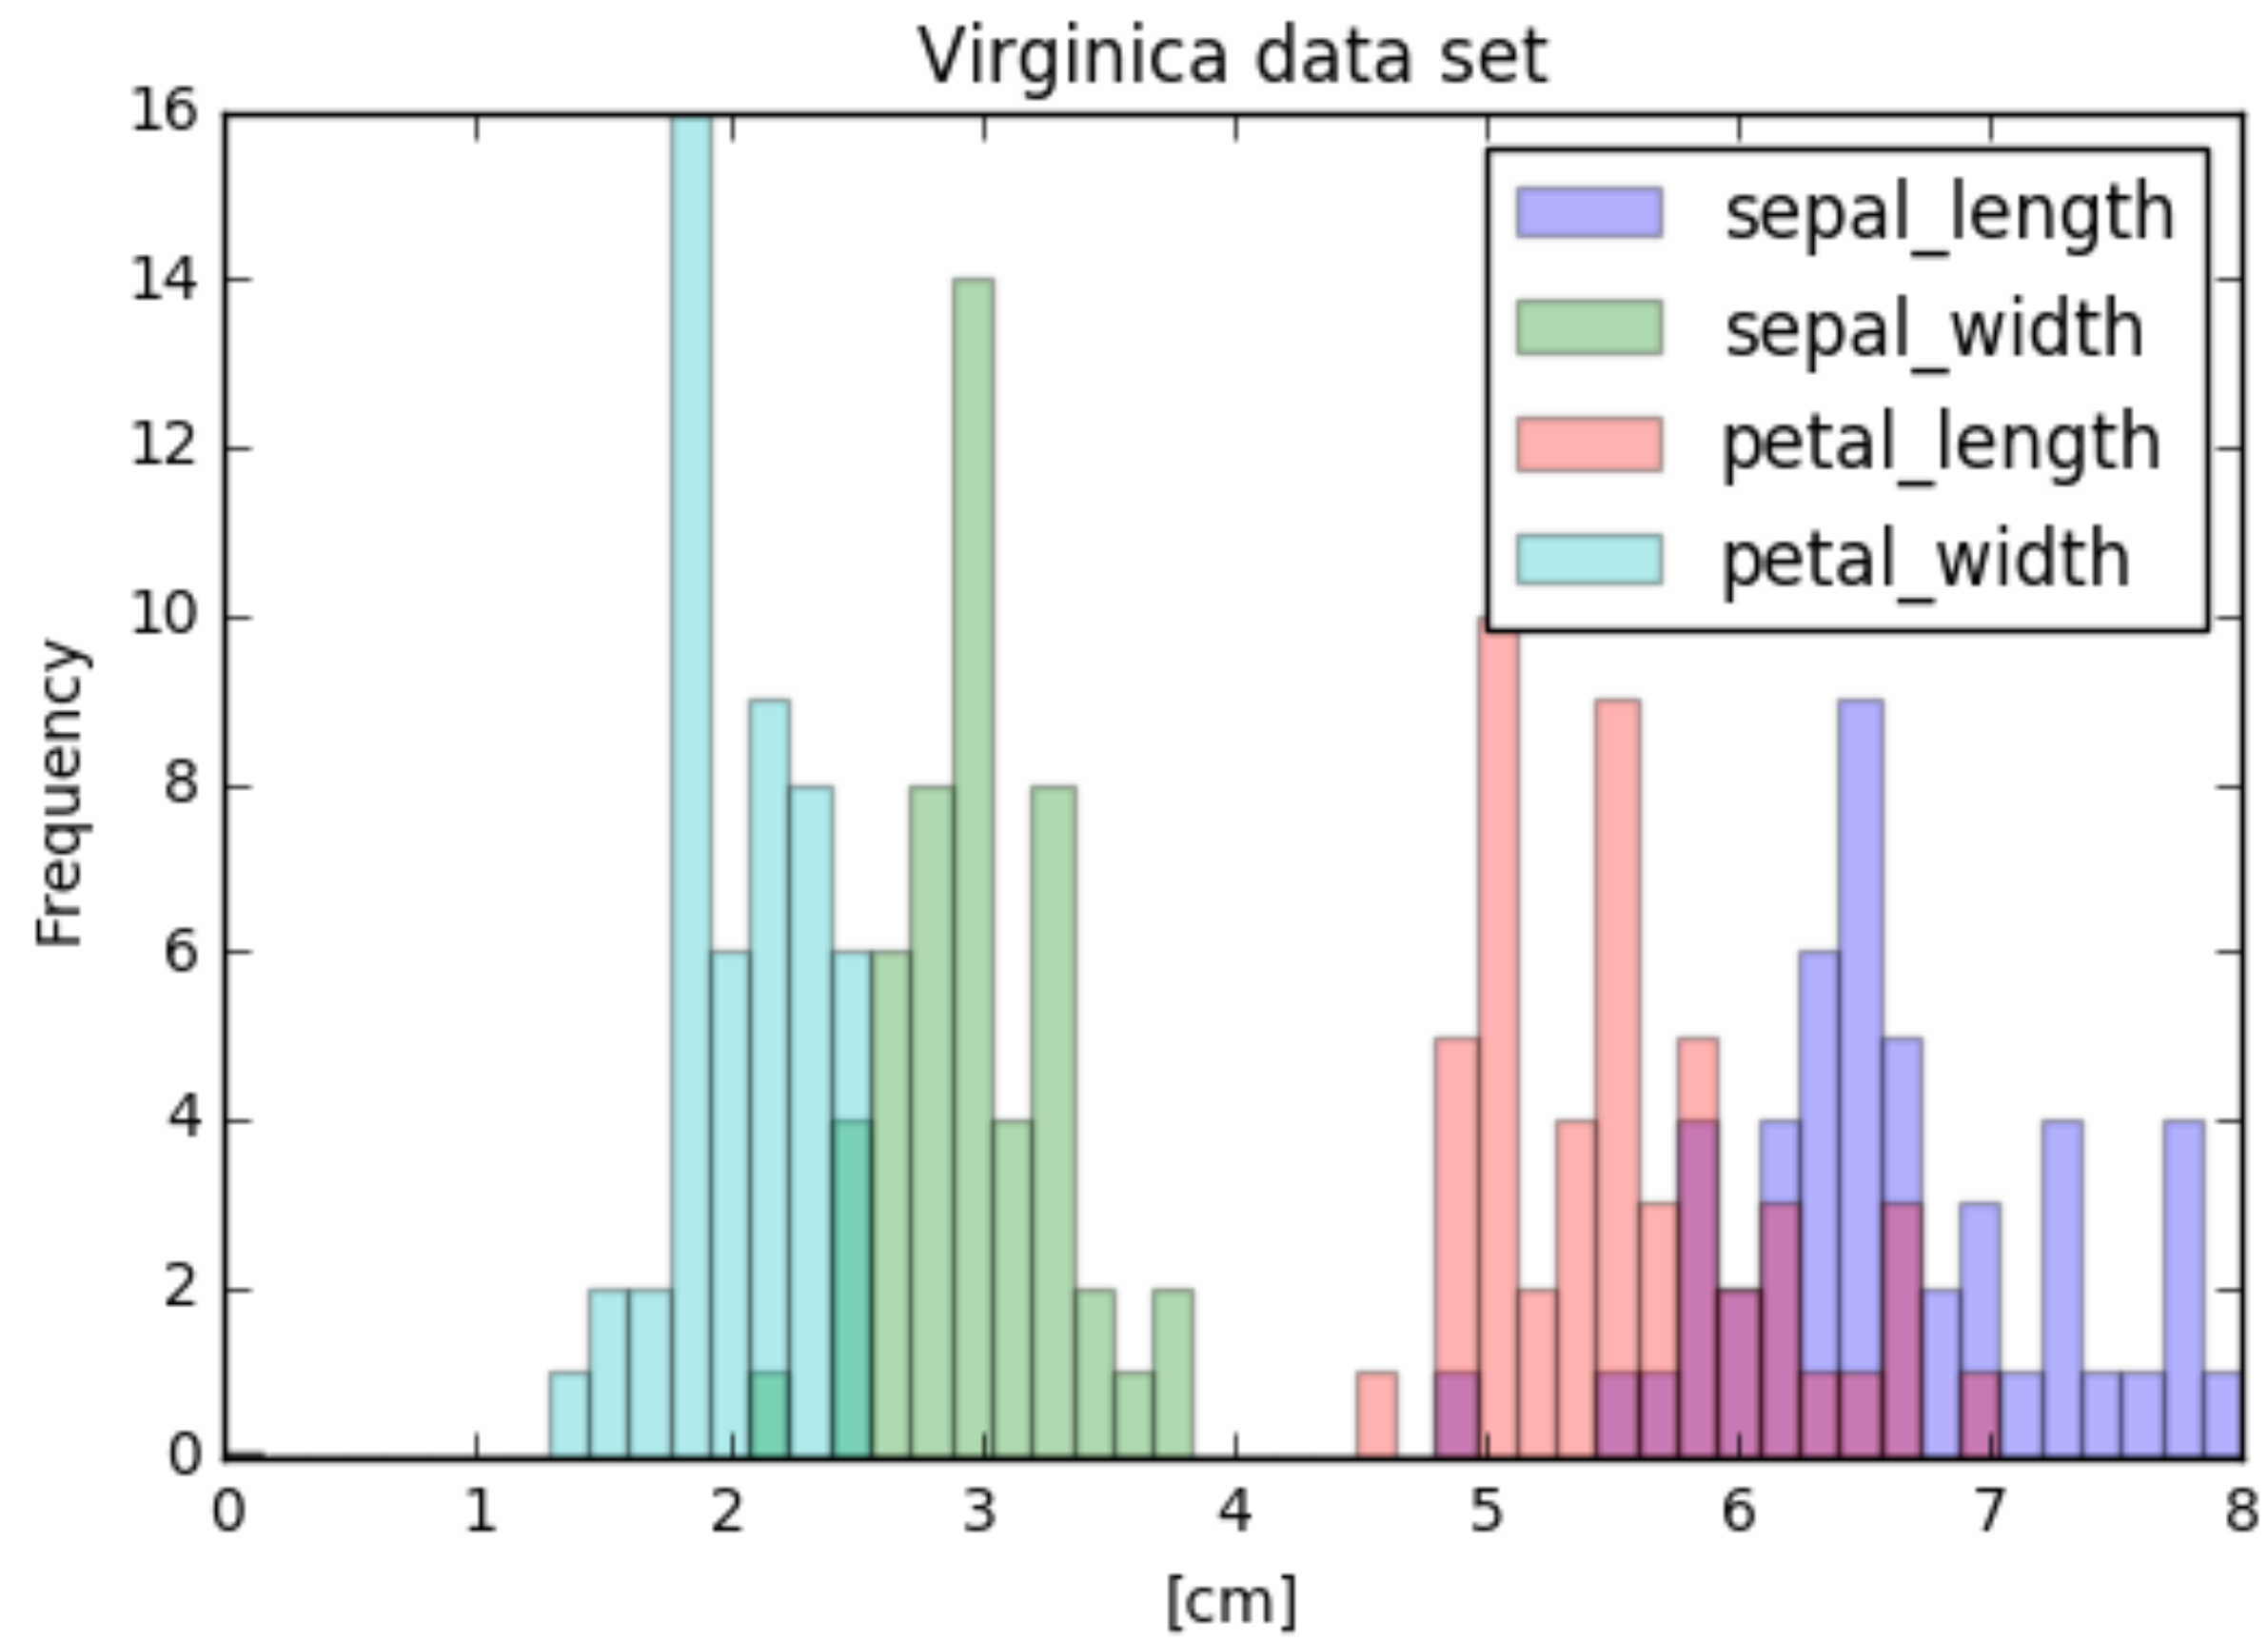
Visual EDA: Setosa data

Visual EDA: Versicolor data

Visual EDA: Virginica data

Statistical EDA: describe()

```
In [31]: describe_all = iris.describe()

In [32]: print(describe_all)
Out[32]:
         sepal_length    sepal_width    petal_length    petal_width
count      150.000000     150.000000      150.000000     150.000000
mean         5.843333       3.057333        3.758000       1.199333
std          0.828066       0.435866        1.765298       0.762238
min          4.300000       2.000000        1.000000       0.100000
25%          5.100000       2.800000        1.600000       0.300000
50%          5.800000       3.000000        4.350000       1.300000
75%          6.400000       3.300000        5.100000       1.800000
max          7.900000       4.400000        6.900000       2.500000

In [33]: describe_setosa = setosa.describe()

In [34]: describe_versicolor = versicolor.describe()

In [35]: describe_virginica = virginica.describe()
```

Computing errors

```
In [36]: error_setosa = 100 * np.abs(describe_setosa -
    ...: describe_all)

In [37]: error_setosa = error_setosa/describe_setosa

In [38]: error_versicolor = 100 * np.abs(describe_versicolor -
    ...: describe_all)

In [39]: error_versicolor = error_versicolor/describe_versicolor

In [40]: error_virginica = 100 * np.abs(describe_virginica -
    ...: describe_all)

In [41]: error_virginica = error_verginica/describe_virginica
```

Viewing errors

```
In [42]: print(error_setosa)
        sepal_length  sepal_width  petal_length  petal_width
count    200.000000    200.000000    200.000000   200.000000
mean      16.726595     10.812913    157.045144   387.533875
std      134.919250     14.984768    916.502136   623.284534
min        0.000000     13.043478      0.000000     0.000000
25%        6.250000     12.500000     14.285714    50.000000
50%       16.000000     11.764706    190.000000   550.000000
75%       23.076923     10.204082    223.809524   500.000000
max       36.206897      0.000000    263.157895   316.666667
```

PANDAS FOUNDATIONS

Let's practice!